

Hardware-assisted Run-time Protection

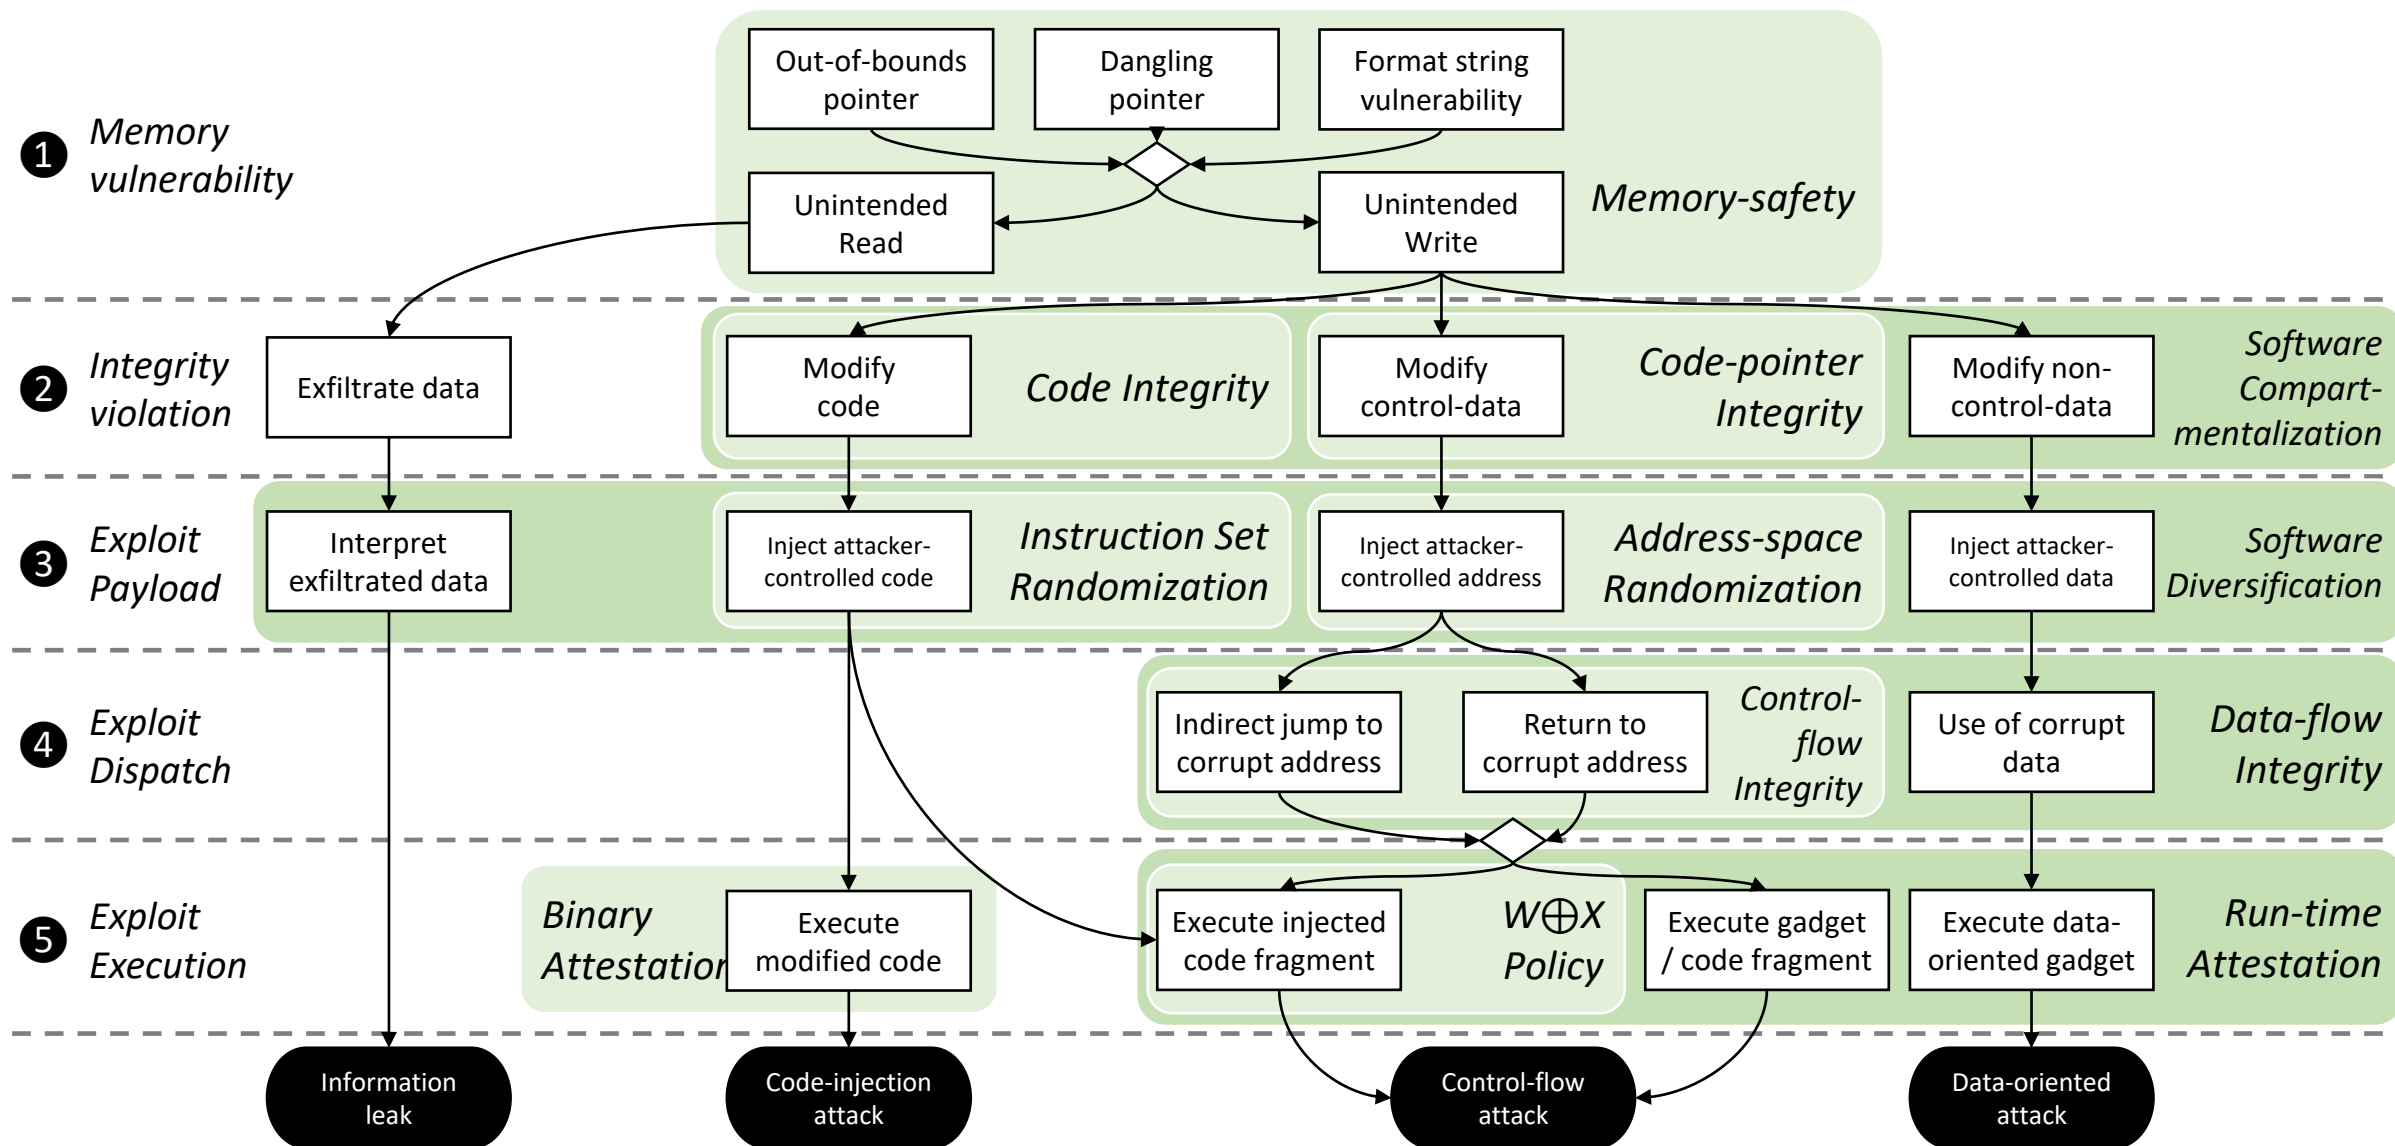
Thomas Nyman[‡], N. Asokan^{†‡}

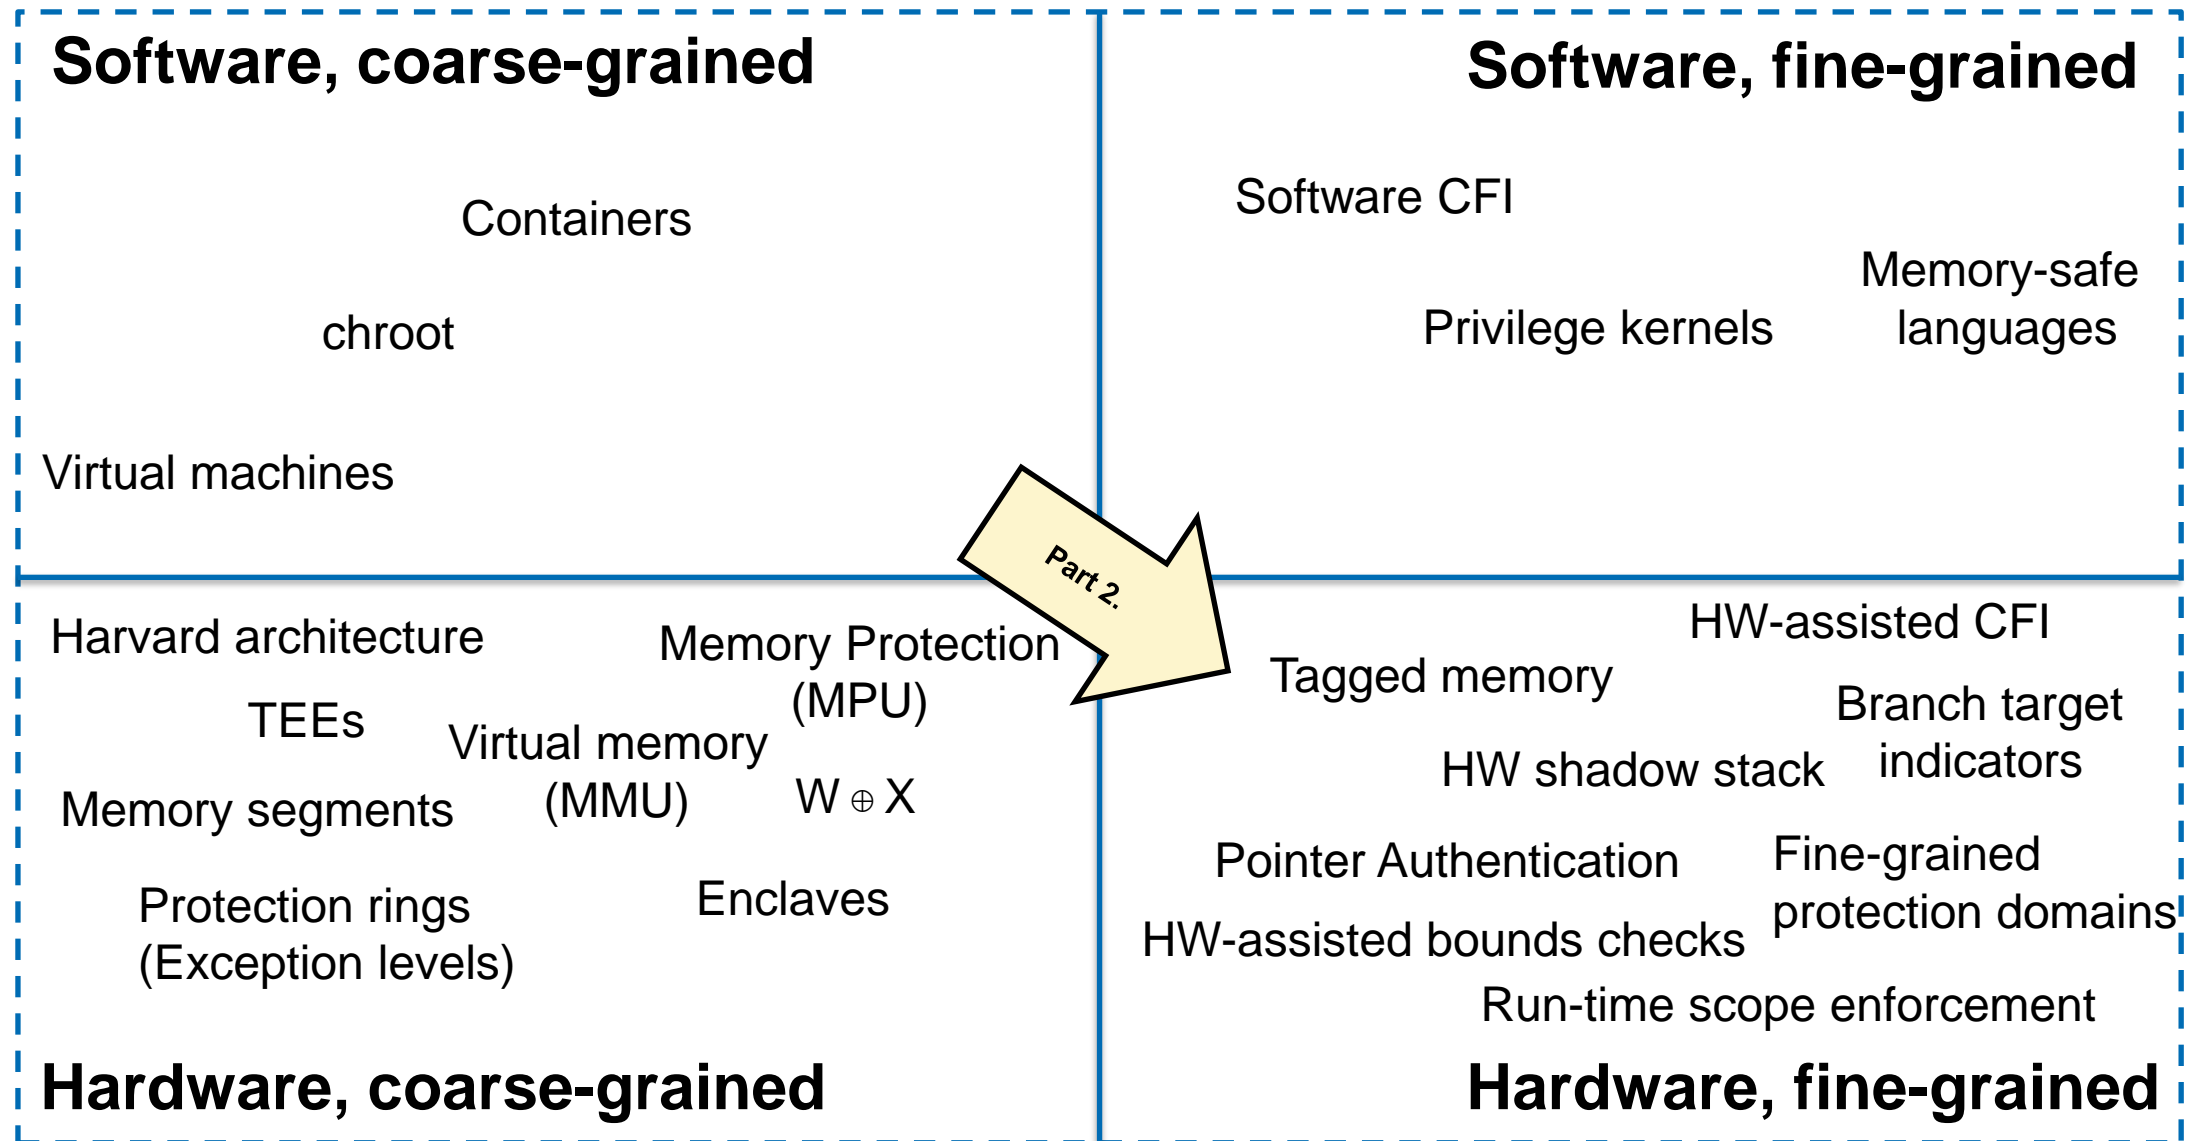
 <https://asokan.org/asokan/>
 @nasokan

Acknowledgements: Hans Liljestrand[†], Lachlan J Gunn[‡], Jan-Erik Ekberg^{‡, §}

[†]) University of Waterloo, [‡]) Aalto University, [§]) Huawei Technologies

Taxonomy of Defenses





Hardware-assisted defenses

How to thwart run-time attacks?

Run-time attacks are now routine

Software defenses incur security vs. cost tradeoffs

Hardware-assisted defenses are attractive

Hardware assisted defenses in CotS processors

ARMv8-A mechanisms

**Pointer Authentication
(PA)**

**Memory Tagging
Extension (MTE)**

**Branch Target
Identification (BTI)**

Intel x84_64 mechanisms

**Memory Protection
eXtension (MPX)**

**Memory Protection Keys
(PKU)**

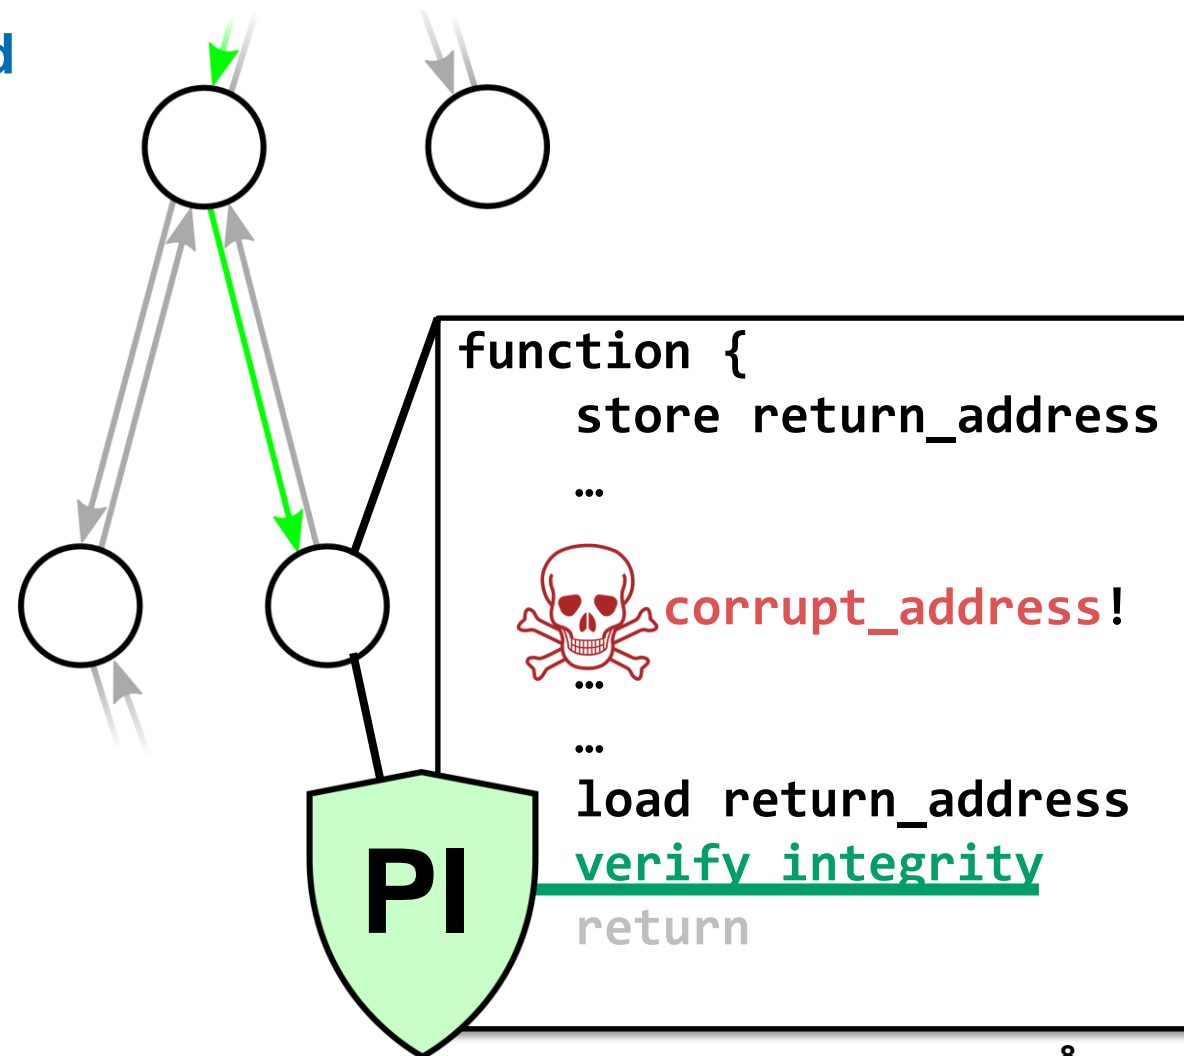
**Control-flow Enforcement
Technology (CET)**

ARMv8-A mechanisms

Pointer Integrity: memory safety for pointers

Ensure **pointers** in memory remain **unchanged**

- **Code pointer integrity implies CFI**
 - Control-flow attacks manipulate code pointers
- **Data pointer integrity**
 - Reduces data-only attack surface



ARMv8.3-A Pointer Authentication



General purpose hardware primitive approximating pointer integrity

- Ensure **pointers** in memory remain **unchanged**

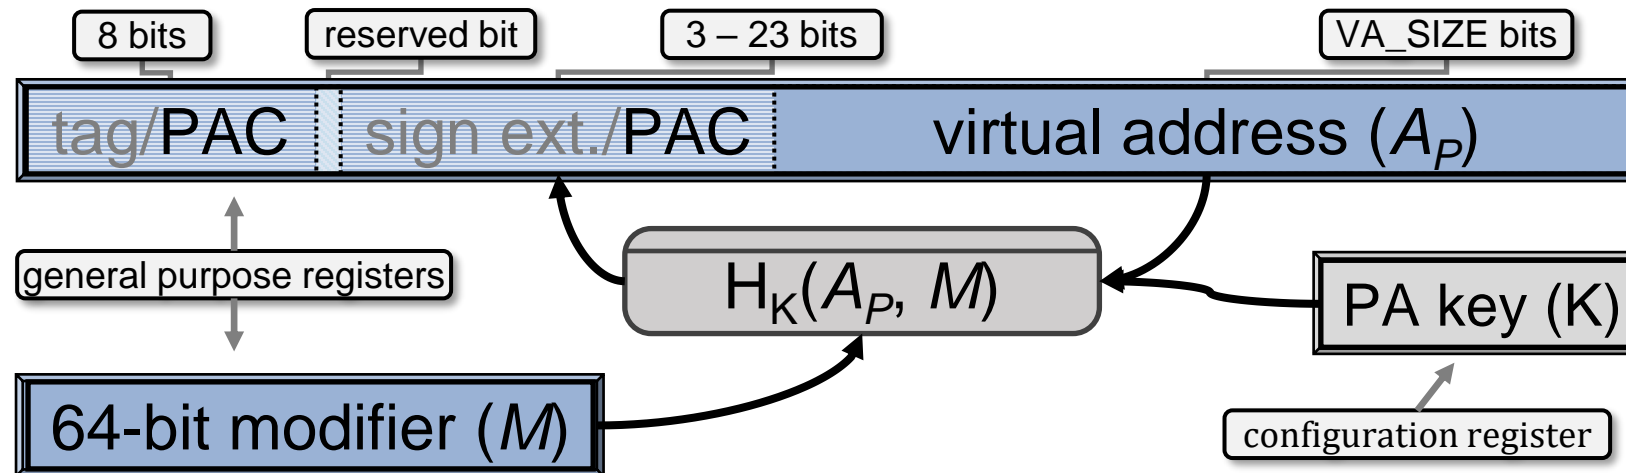
Introduced in ARMv8.3-A specification (2016) to be improved in ARM-8.6-A (2020)

- First compatible processors 2018 (Apple A12 / [iOS12](#))
- Support in [Linux 5.0](#)
- Instrumentation support in [GCC 7.0](#) ([-msign-return address](#), deprecated in [GCC 9.0](#) [-mbranch-protection=pac-ret\[+leaf\]](#) GCC 9.0 and newer)

ARMv8.3-A PA – PAC Generation

Adds Pointer Authentication Code (**PAC**) into unused bits of pointer

- Keyed, tweakable **MAC** from **pointer address** and 64-bit **modifier**
- PA keys protected by hardware, modifier decided where pointer **created and used**



ARMv8.3-A PA – Key management and instructions

Keys for PAC generation and verification

APIAKey_EL1	Key A for instruction address PACs
APIBKey_EL1	Key B for instruction address PACs
APDAKey_EL1	Key A for data address PACs
APDBKey_EL1	Key B for data address PACs
APGAKey_EL1	Key for generic authentication

PA Instructions

PAC<i d><a/b> <Xd> <Xm>	Add PAC to address in <i>Xd</i> using modifier in <i>Xm</i>
AUT<i d><a/b> <Xd> <Xm>	Authenticate address in <i>Xd</i> using modifier in <i>Xm</i>
PACGA <Xd> <Xn> <Xm>	Calculate generic PAC for data in <i>Xn</i> using modifier in <i>Xm</i>
XPAC<i d> <Xd>	Strip PAC for address in <i>Xd</i>
BRA<a b> <Xn> <Xm>	Branch to address in <i>Xn</i> after authenticating it with modifier in <i>Xm</i>
BLRA<a b> <Xn> <Xm>	As BRA but perform branch with link
RETA<a b>	Authenticate address in LR with SP as modifier and return
ERETA<a b>	Authenticate address in ELR with SP as modifier and exception return
LDRA<a b> <Xt> <Xn>	Authenticate address in <i>Xn</i> using modifier zero and load value to <i>Xt</i>

operate on instruction keys only

operate on data keys only

PA-based protection schemes

PA instructions are **primitives**, assembled to form **protection schemes**

Two main components:

- When are pointers “PACed” and “unPACed”?
- Which modifier is used at a given point?

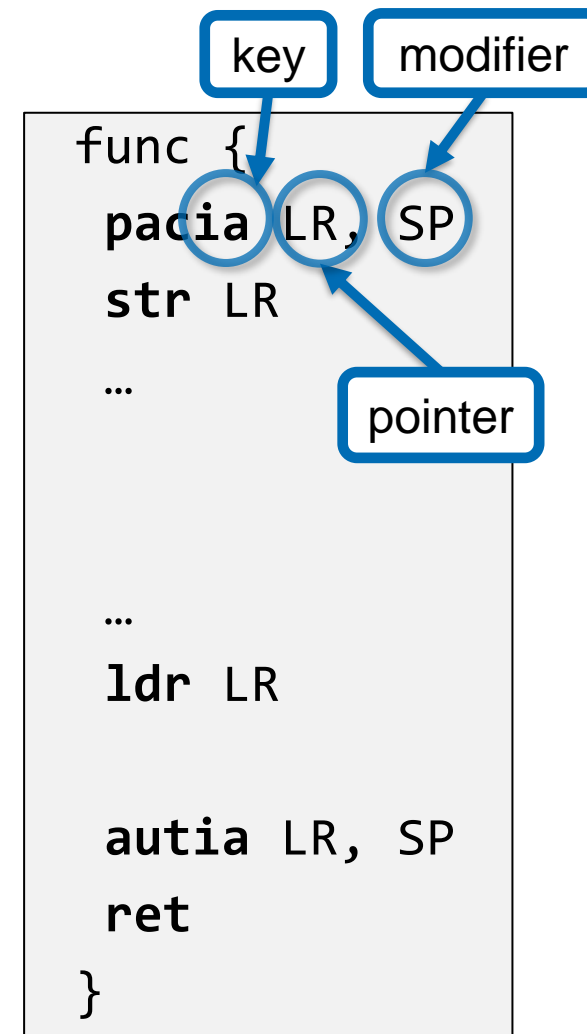
What should the modifier be for a given pointer?

- For **security**: using many different modifiers makes **replay attacks harder**
- For **functionality**: large numbers of modifiers are **hard to keep track of**

PA prevents arbitrary pointer injection

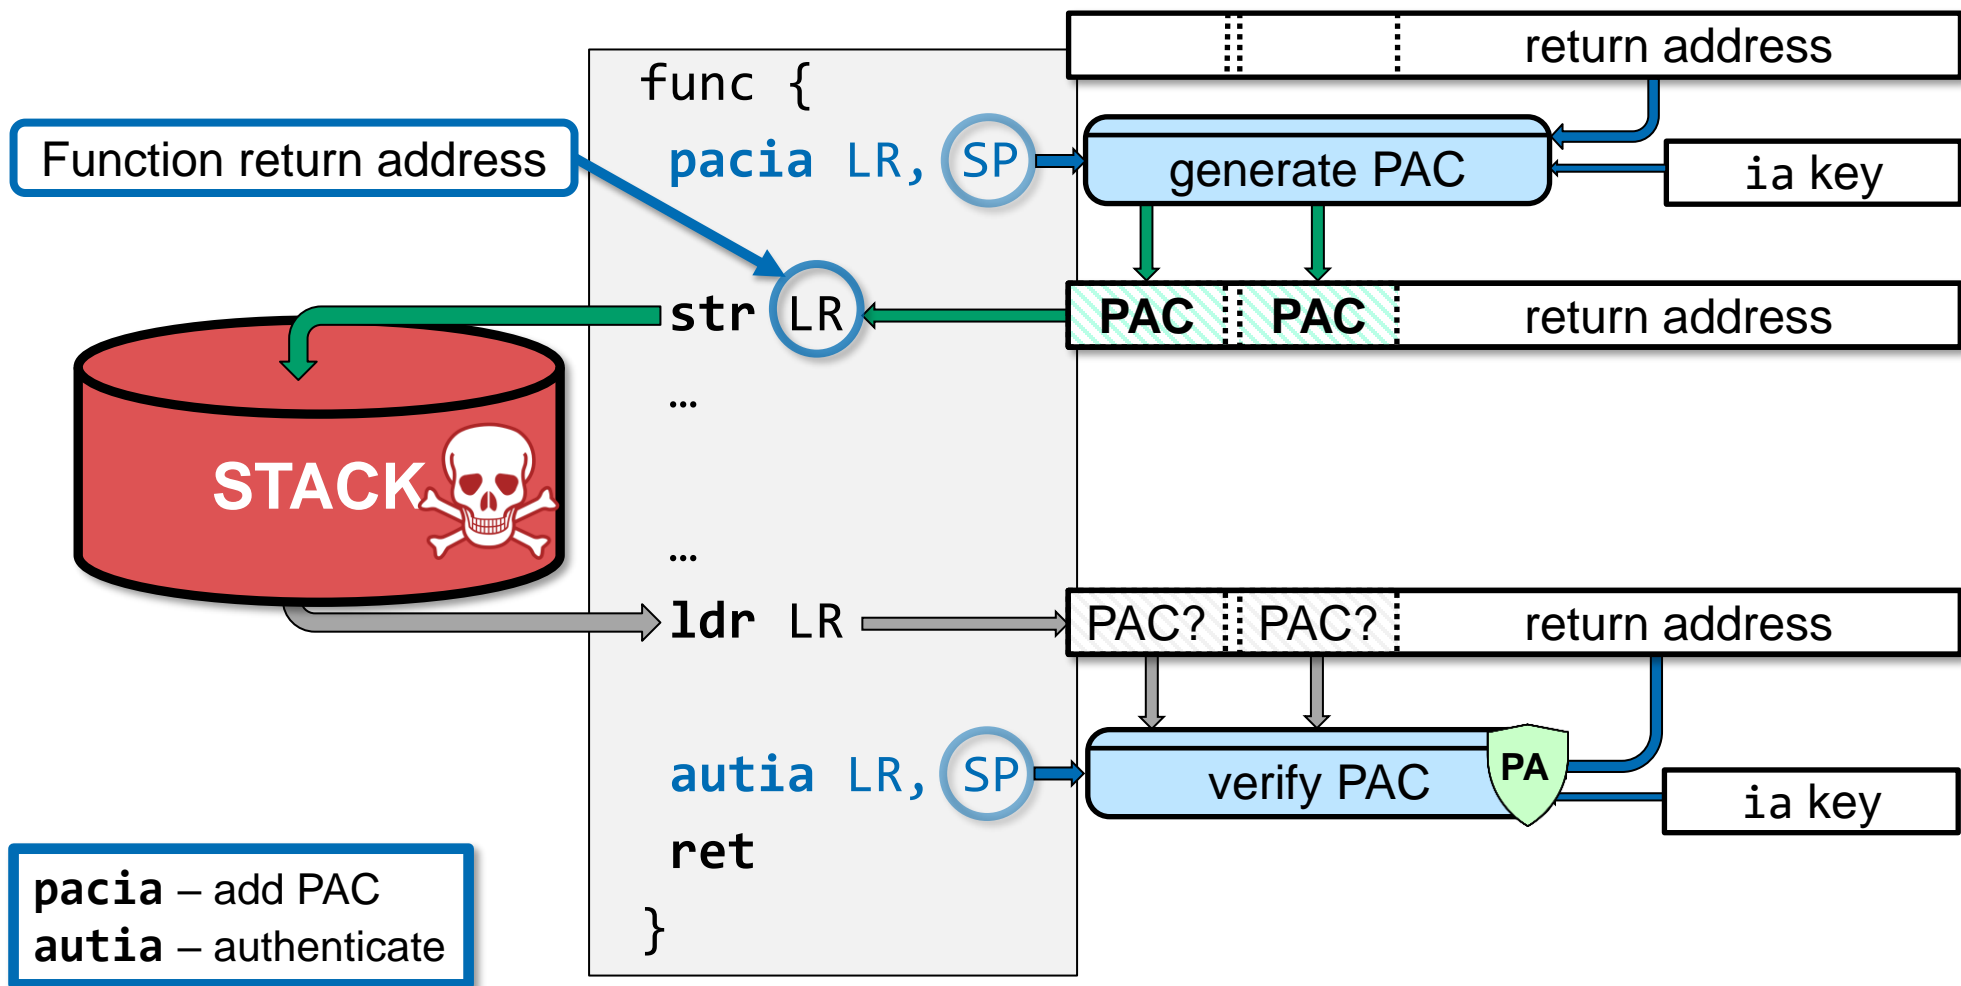
- **Modifiers do not need to be confidential**
 - Visible or inferable from the code section / binary
- **Keys are protected by hardware and set by kernel**
 - Attacker cannot generate PACs

pacia – add PAC
autia – authenticate



Example: -msign-return-address

Deployed in GCC 5.0 and LLVM/Clang 7.0



PA return address protection as a canary

The signed return address effectively **is a canary**:

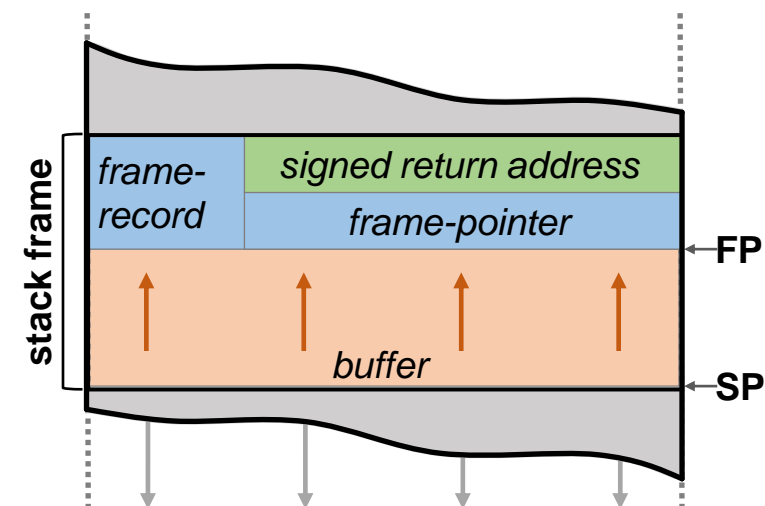
- Any overflow that corrupts the return address is detected

More powerful than -stack-protector canaries:

- Does **not require reference value**
- Can be **bound to contextual information** (e.g., the SP value)
- Protects return address against arbitrary writes

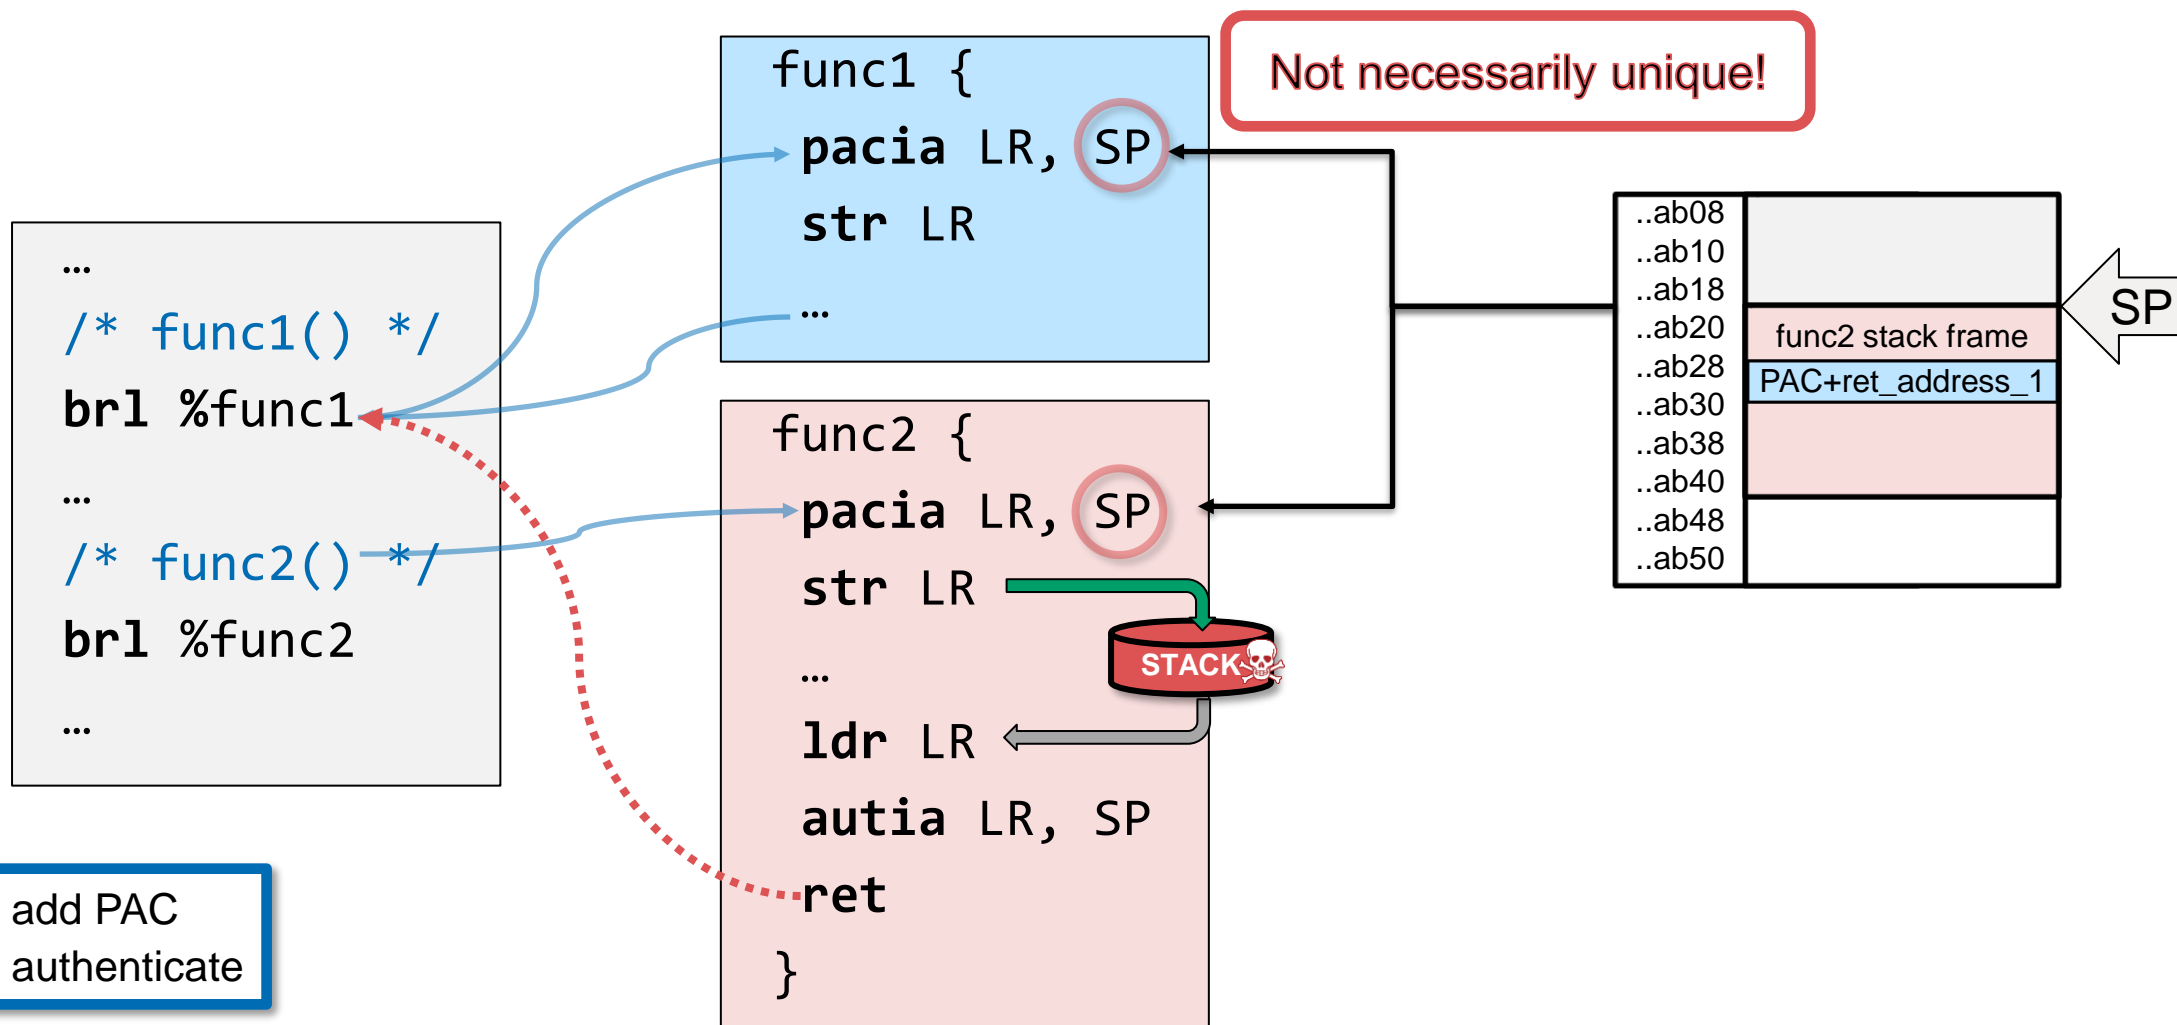
Also has similar weaknesses:

- Existing return addresses can be **reused**



PA only approximates fully-precise pointer integrity

Adversary may **reuse** PACs



PARTS

Modifier: based on pointer type

- Assigned at compile-time based on C type
- “this pointer really points to this type of data or function”

On-load, on-branch authentication

- Branching with combined auth+branch instruction (**lbraa**)
- Iterating an array uses **only one authentication**

pacda	– add PAC with data A-key
autda	– authenticate
pacia	– add PAC with instr A-key
lbraa	– authenticate and branch

```
// *ptr
...
ldr Xptr, <memory>
mov Xmod, #type_id
autda Xptr, Xmod
<something> [Xptr]
```

```
// ptr = ...
...
mov Xmod, #type_id
pacia Xptr, Xmod
```

PACed only on pointer creation!

```
// ptr();
...
...
mov Xmod, #type_id
lbraa Xptr, Xmod
...
```

Authenticated on use

ARMv8.5-A Memory Tagging Extension



Ensures memory accesses are safe by comparing tag in pointer with tag in memory

Introduced in ARMv8.5-A specification (announced September 2018)

- Support in Linux [proposed July 2019](#)
- Stack Tagging will become available in [LLVM 9](#)
- Heap Tagging support planned

ARMv8.5-A MTE

Address tags stored in top byte of pointer

- uses existing top-byte ignore feature

Allocation tags stored by transparently by hardware and cached

- 4-bit tag per 16-byte granule of memory

Mismatch between tags reported either:

- synchronously (precise check during testing), or
- asynchronously (imprecise checks after deployment)

ARMv8.5-A MTE

```
char *ptr = new char [16]; // memory colored
```



```
ptr[17] = 42; // color mismatch -> overflow
```

```
delete [] ptr; // memory re-coloured on free
```



```
ptr[10] = 10; // color mismatch -> use-after-free
```

ARMv8.5-A MTE – Instructions

MTE Instructions

IRG	Insert random address tag to address in register
GMI	Manipulate excluded set of tags for IRG
ADDG/SUBG	Arithmetic on addresses with tags for creating pointer to objects on stack
SUBP(S)	56-bit subtract allowing address tag in top byte of pointers to be ignored
LDG/STG/STZG	Get or set allocation tags for granule (STGZ also initializes data to zero)
ST2G/STZ2G	Like STG or STZG but operate on two granules of memory at a time
STGP	Store both tag and data to memory
LDGM/STGM/STZGM	Bulk tag manipulation for initializing or serializing tags by system software

ARM. [Arm v8.5-A Memory Tagging Extension](#), whitepaper 2019

ARM. [Arm® Architecture Reference Manual Armv8, for Armv8-A architecture profile](#). Version E.a. July 2019

ARMv8.5-A Branch Target Identification



Hardware-assisted CFI similar to Intel CET Indirect Branch Tracking

Introduced in ARMv8.3-A specification (2016)

- Support for Linux proposed [May 2019](#)
- Instrumentation support in [GCC 9.0](#) ([-mbranch-protection=standard|bti](#))

ARMv8.5-A BTI

Indirect branches to **guarded code regions** require **marker instructions**

- compiler places marker potential indirect branch targets
- two classes of targets: **calls** and **jumps** (RET instructions not restricted by BTI)

Branch sources

BTI call type branches

BLR ...	Indirect function calls
BR <x16 x17>	PLT entries and tail calls

BTI jump type branches

BR ... (except x16 x17)	Branches to jump tables
-------------------------	-------------------------

BTI Marker Instructions

BTI < c / j / cj >	Branch Target Identification for c =calls, j =jumps, cj =calls or jumps
BRK	Breakpoint Instruction
HLT	Halting breakpoint
PACIASP / PACIBSP	Create PAC for Instruction address in LR using key A/B and SP as modifier

Taxonomy of Defenses

